



Sonata: Multi-Database Transactions Made Fast and Serializable

Chuzhe Tang¹, Zhaoguo Wang¹, Jinyang Li², Haibo Chen¹

¹Institute of Parallel and Distributed Systems, Shanghai Jiao Tong University

² New York University



SHANGHAI JIAO TONG
UNIVERSITY



Transaction Abstraction

- Transactions are an important programming abstraction.
 - Ensures correctness in the presence of concurrent operation and failures.

**Large-Scale
Applications**



Safeguarding Critical Business Logic

**Transactional
Databases**



Transaction Processing

Transaction Logic

Checking	Savings
\$50	\$50

T1: withdraw(Checking, \$100)

```
Begin Transaction
total := s_bal + c_bal
if total >= 100:
    c_bal -= 100
Commit Transaction
```

T2: withdraw(Savings, \$100)

```
Begin Transaction
total := c_bal + s_bal
if total >= 100:
    s_bal -= 100
Commit Transaction
```

Transaction Processing

Transaction Logic

Submit

Transactional Database

Checking	Savings
\$50	\$50

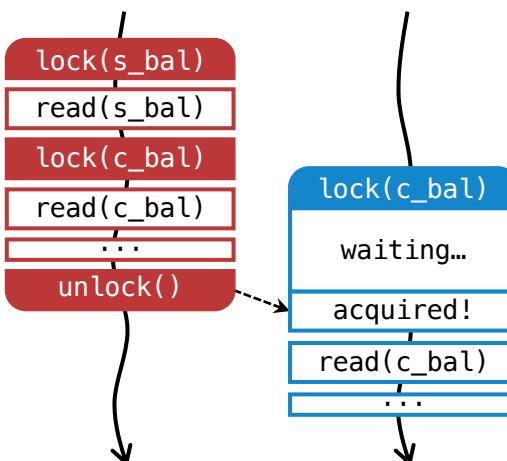
T1: withdraw(Checking, \$100)

```
Begin Transaction
total := s_bal + c_bal
if total >= 100:
    c_bal -= 100
Commit Transaction
```

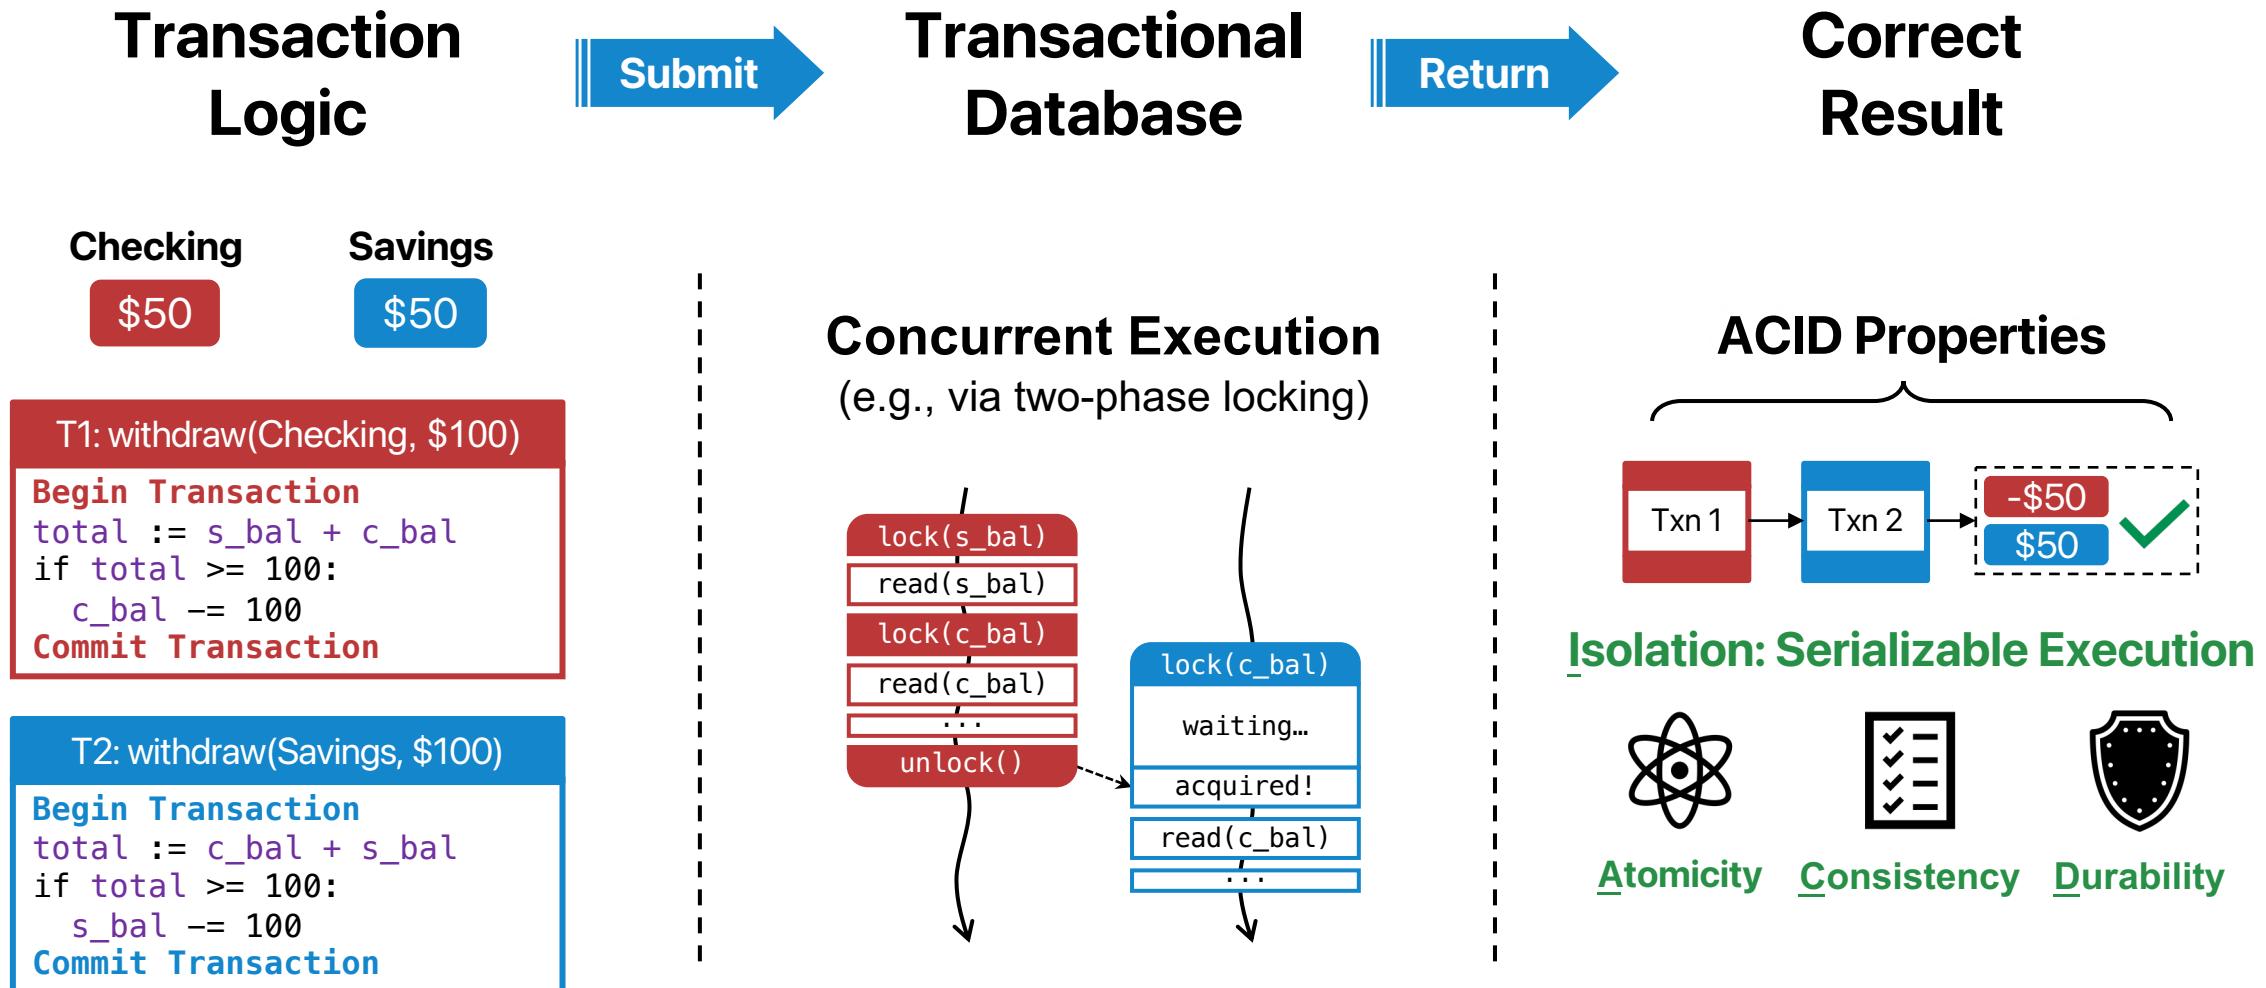
T2: withdraw(Savings, \$100)

```
Begin Transaction
total := c_bal + s_bal
if total >= 100:
    s_bal -= 100
Commit Transaction
```

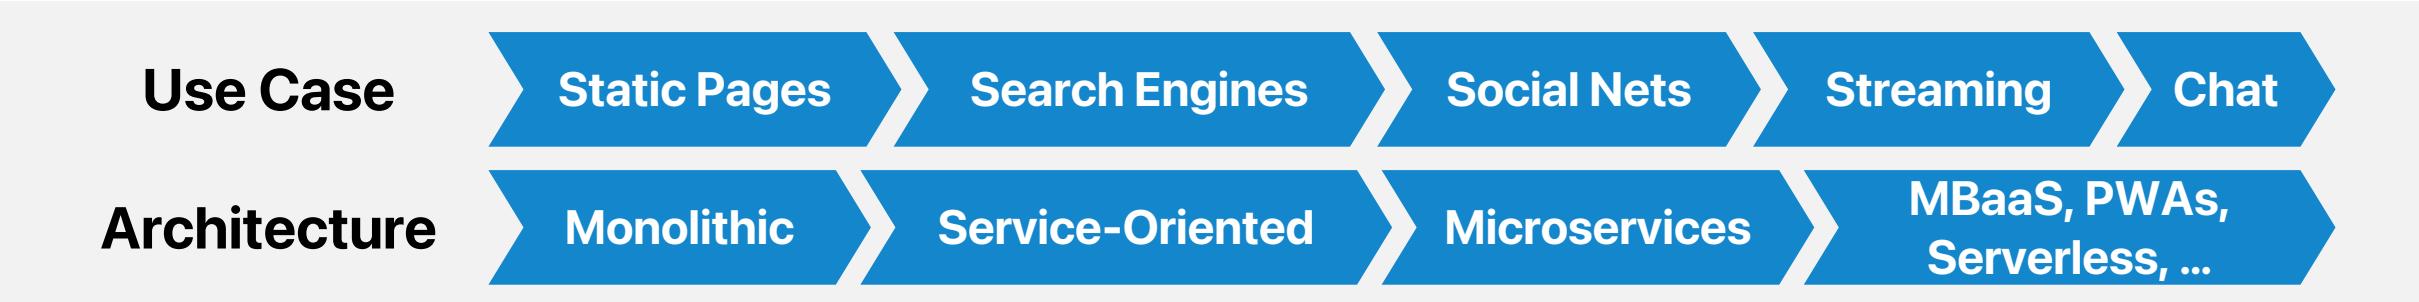
Concurrent Execution (e.g., via two-phase locking)



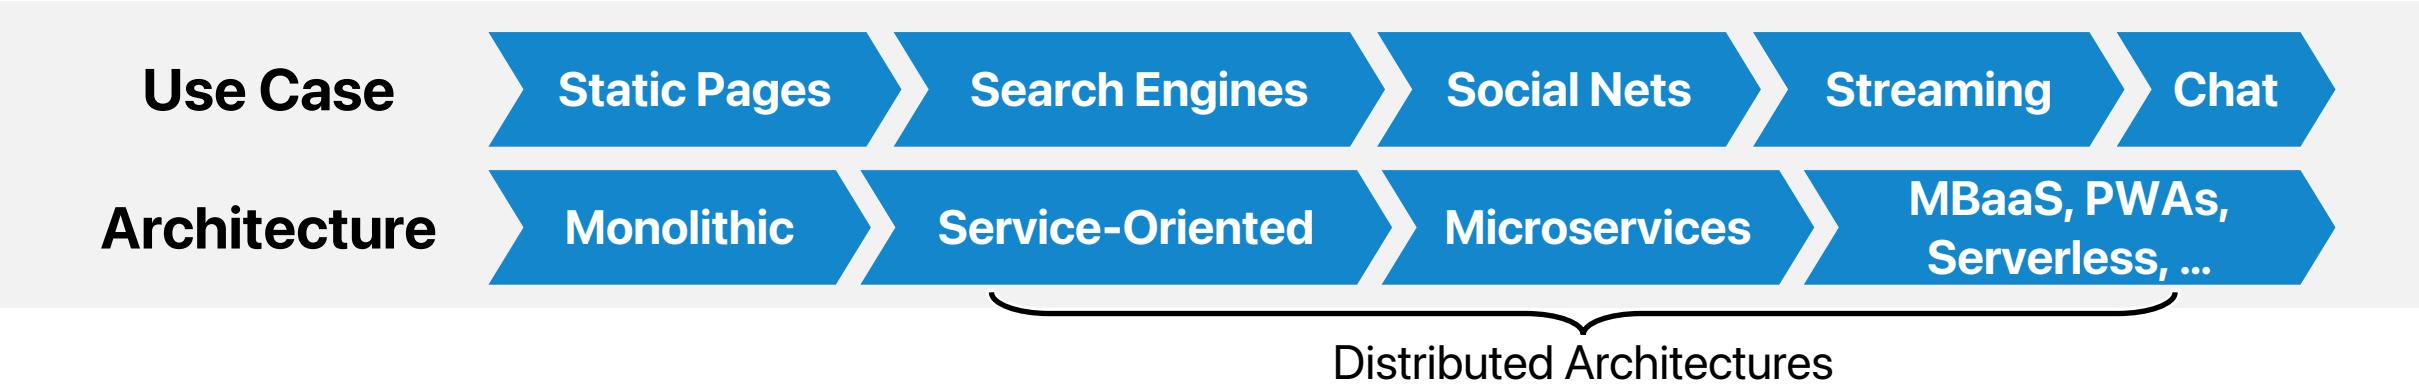
Transaction Processing



Today: Evolving Applications



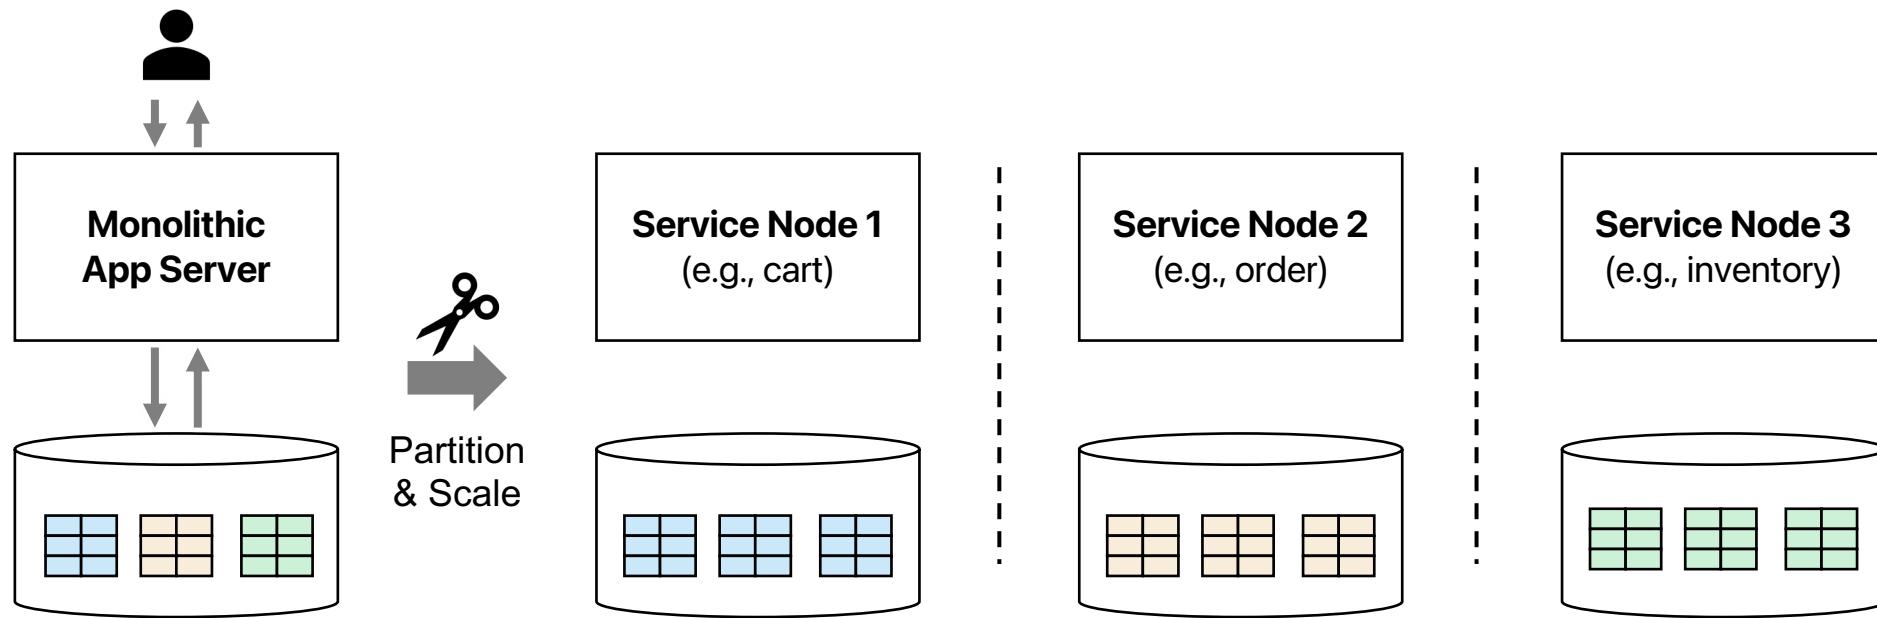
Today: Evolving Applications



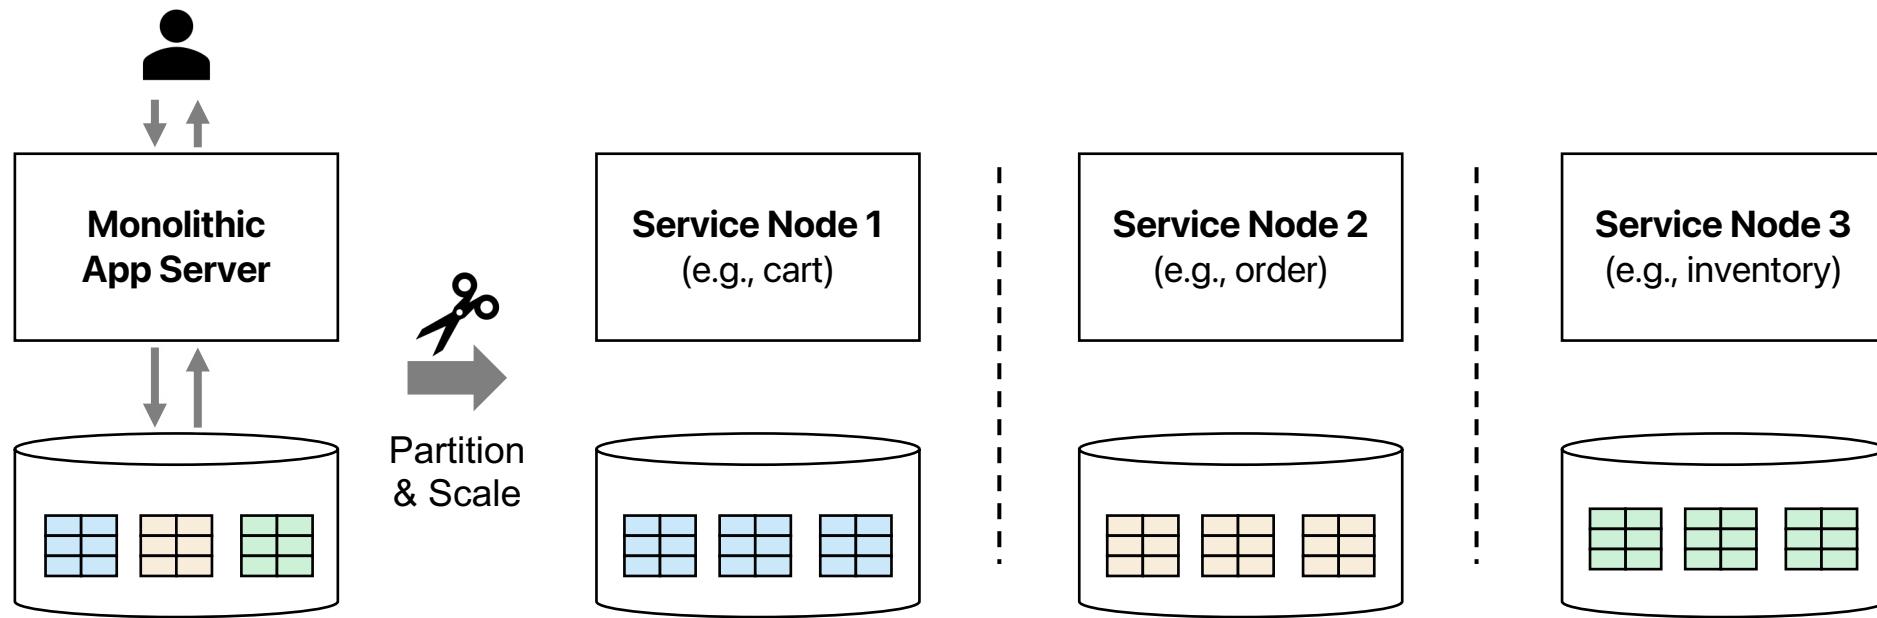
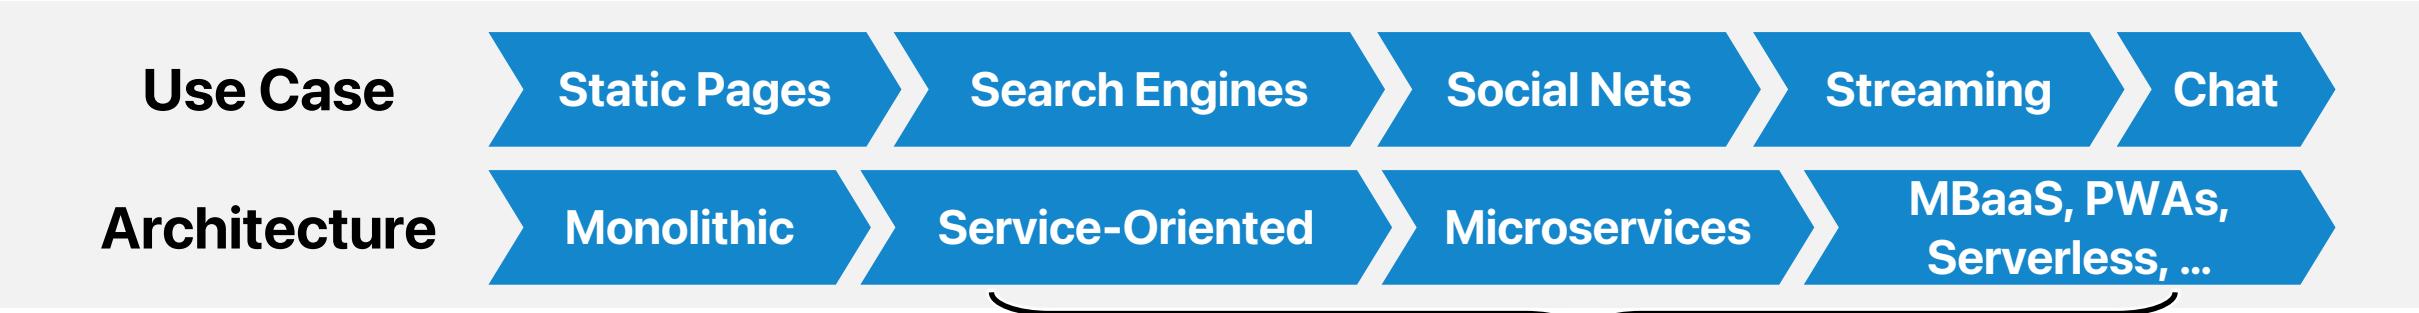
Today: Evolving Applications



Distributed Architectures

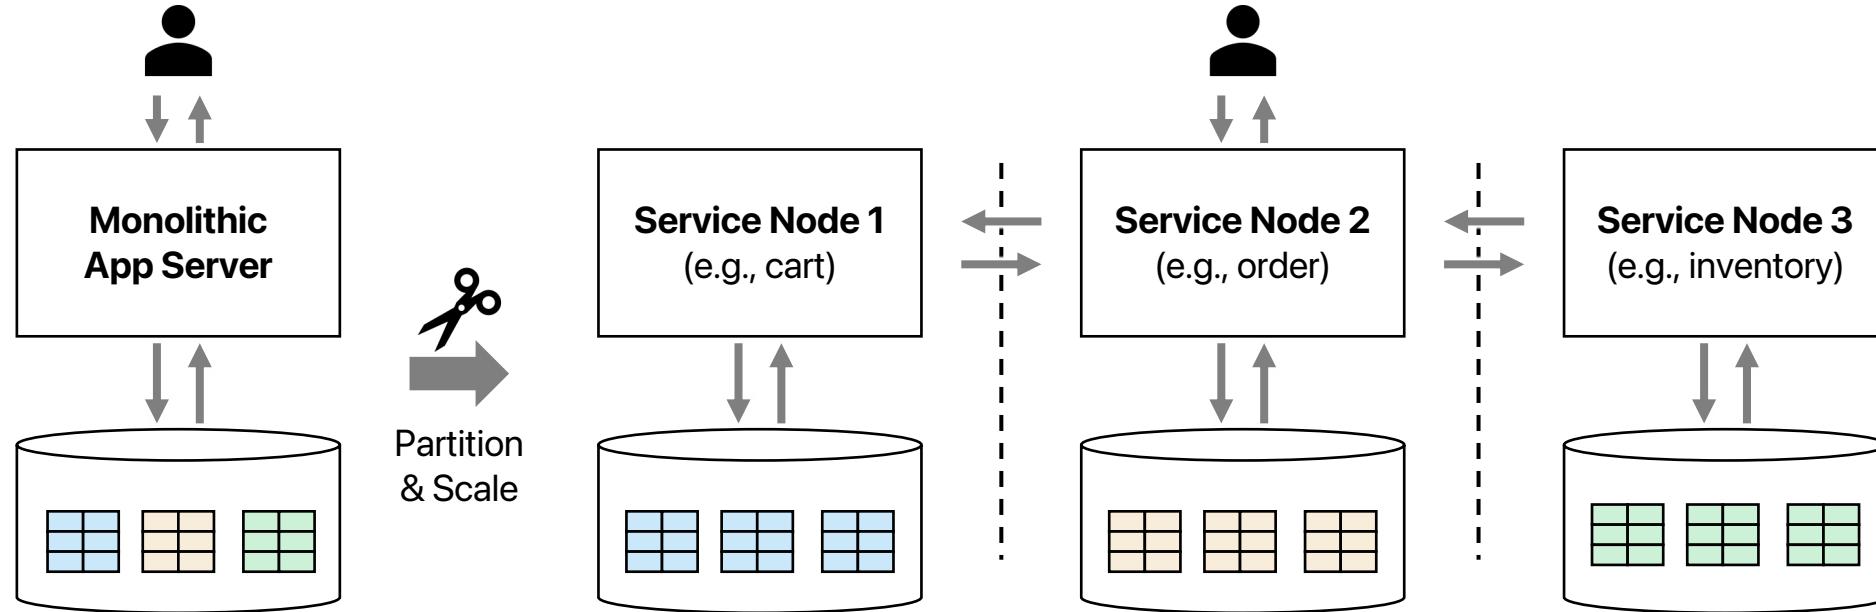


Today: Evolving Applications

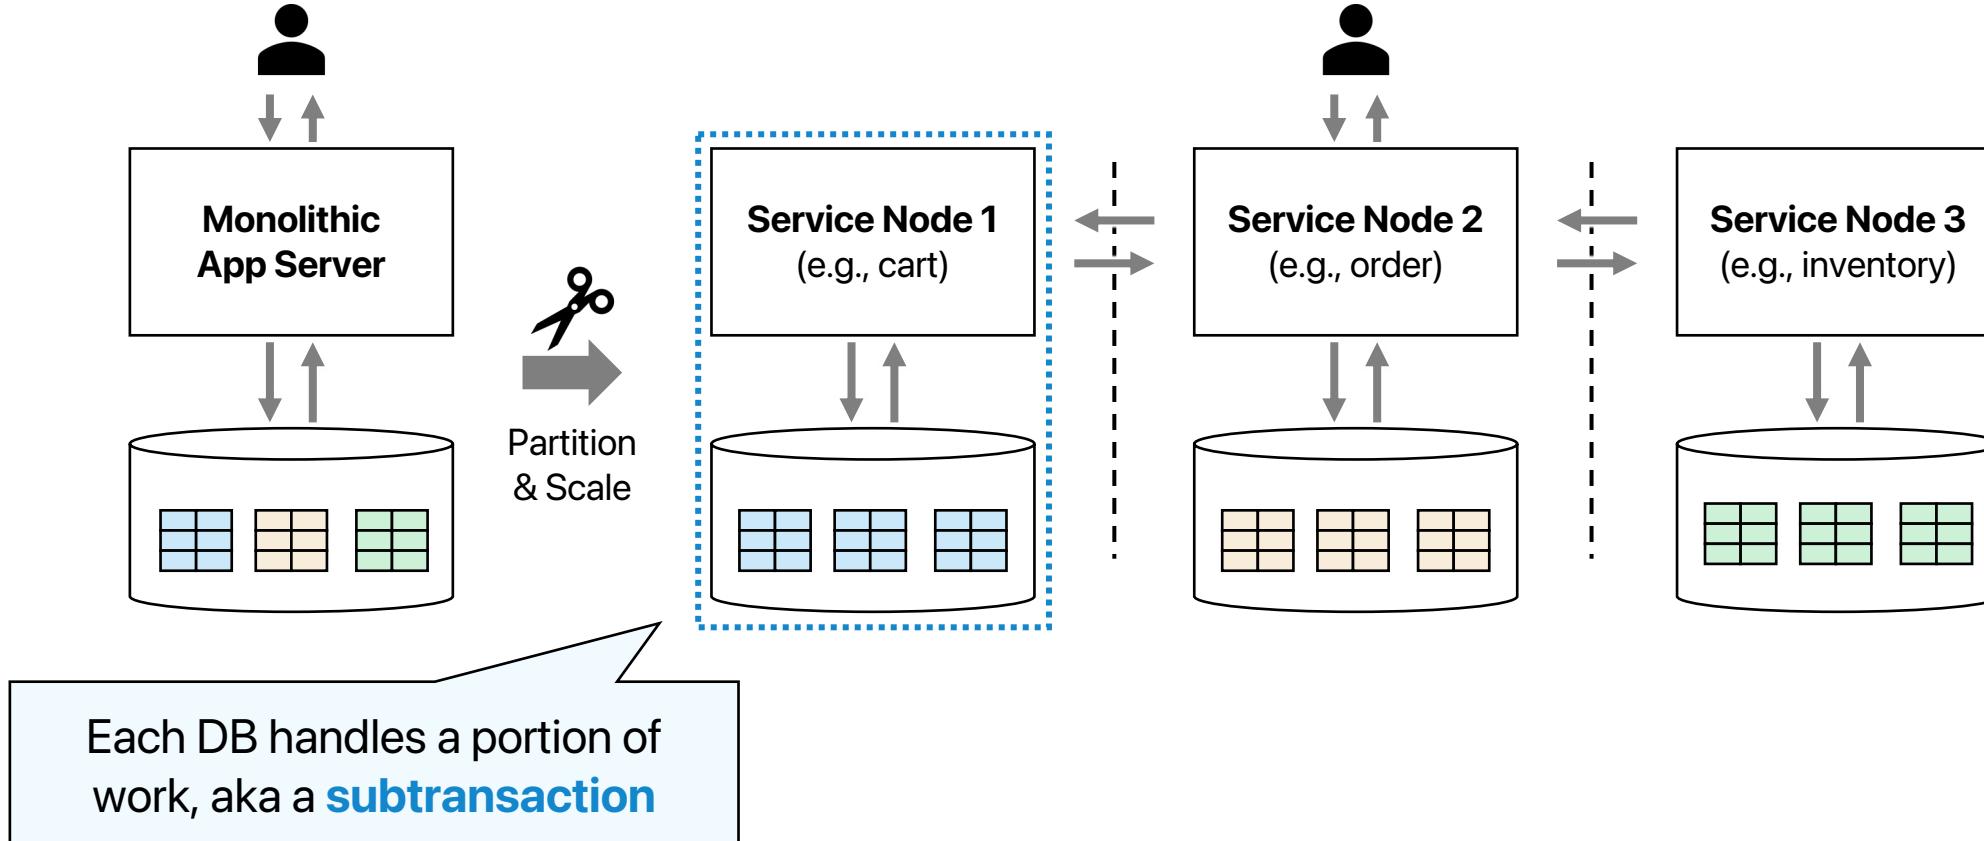


Better scalability in applications'
development, deployment, and maintenance

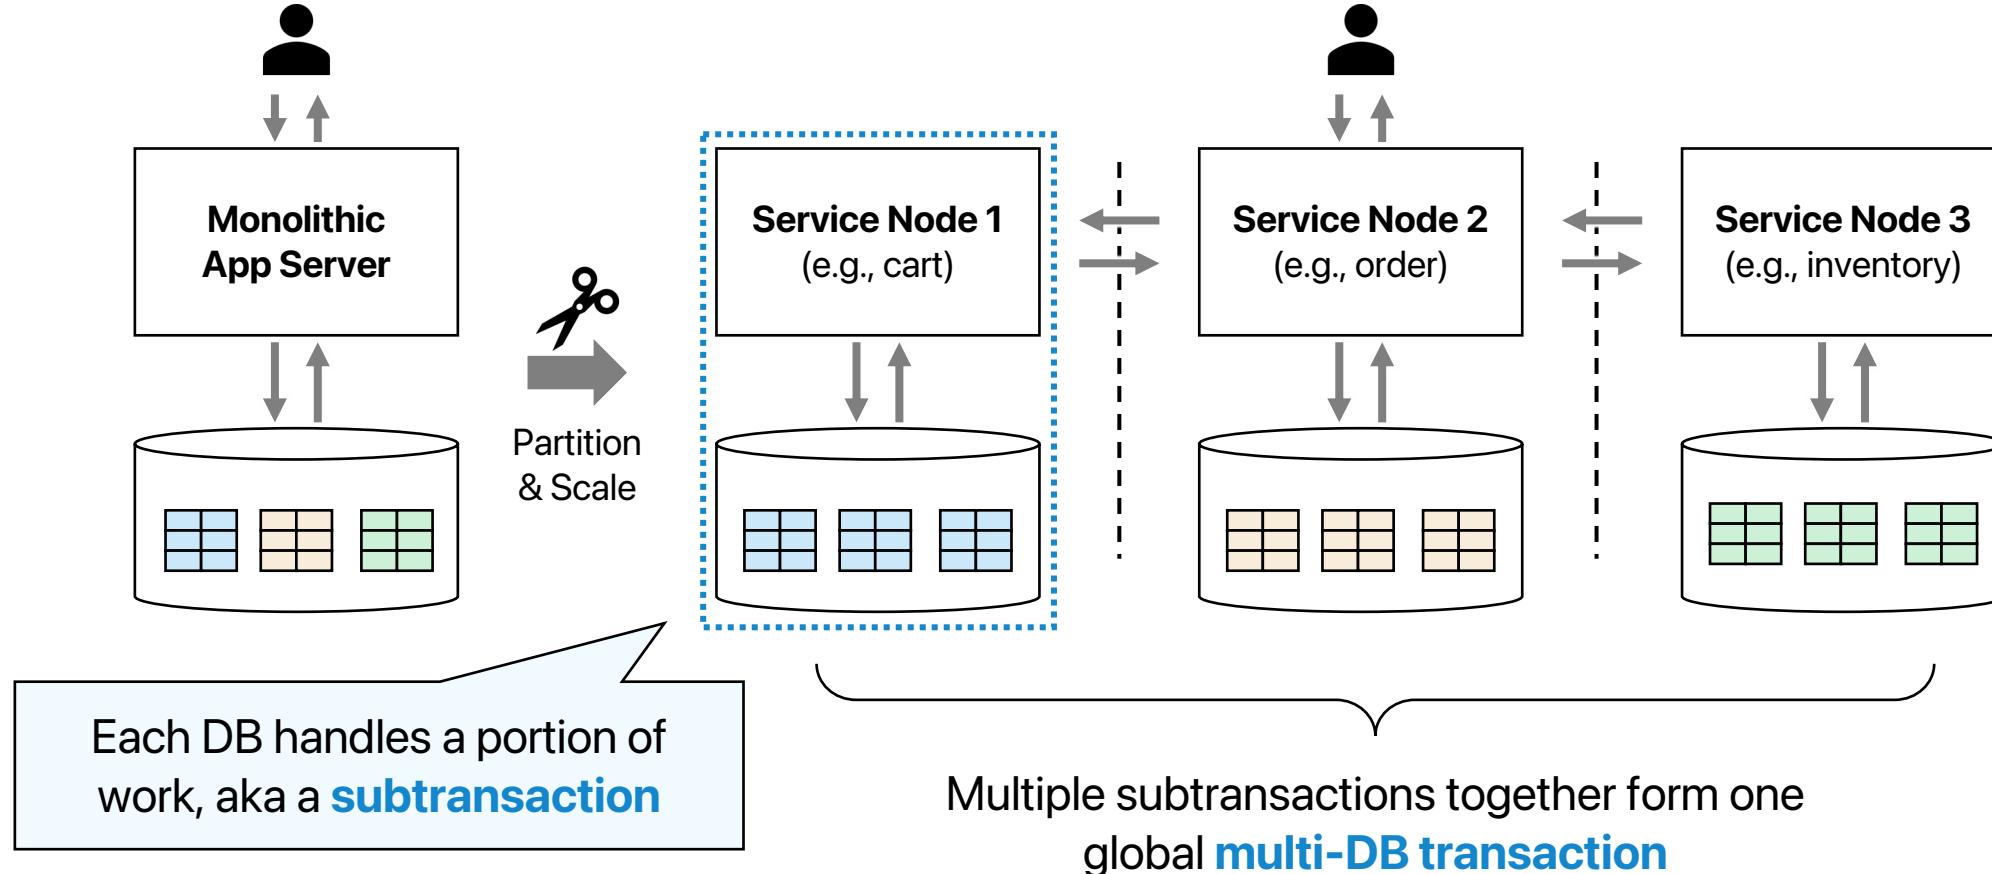
Today: Multi-Database Transactions



Today: Multi-Database Transactions



Today: Multi-Database Transactions



Goal: Global Serializability

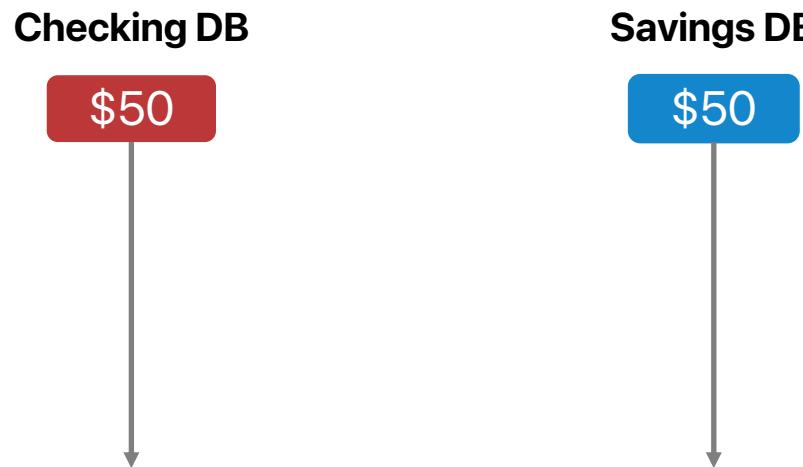
- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.

Goal: Global Serializability

- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.
- **While helpful for correctness, it is non-trivial to realize.**

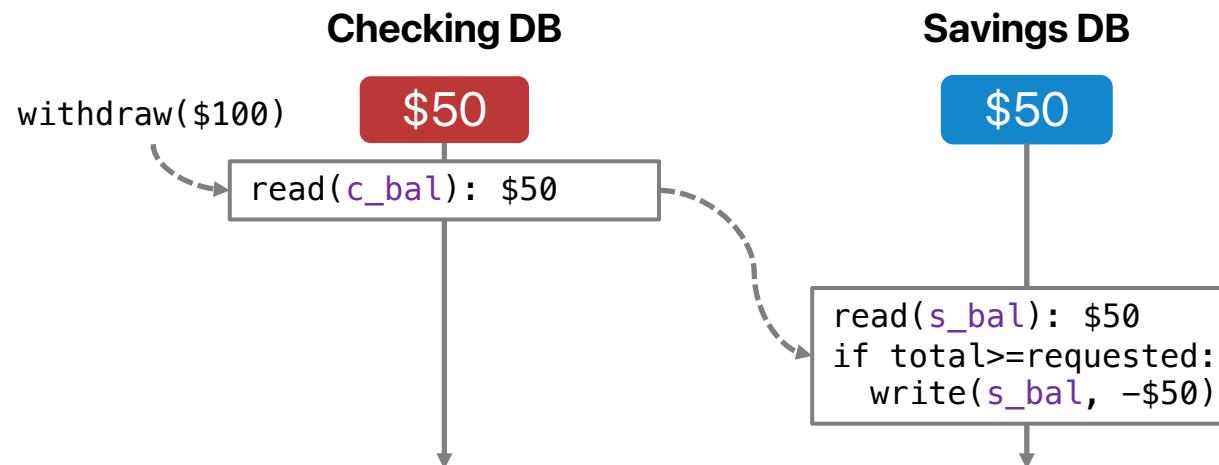
Goal: Global Serializability

- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.
- **While helpful for correctness, it is non-trivial to realize.**
 - Even if all DBs are locally serializable, global serializability can still be violated.



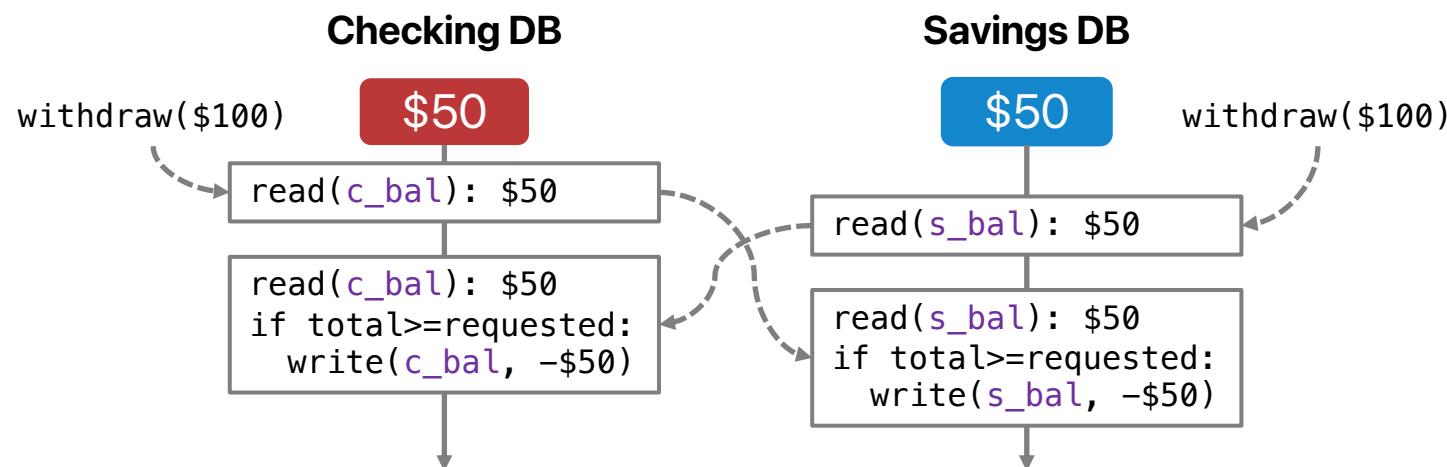
Goal: Global Serializability

- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.
- **While helpful for correctness, it is non-trivial to realize.**
 - Even if all DBs are locally serializable, global serializability can still be violated.



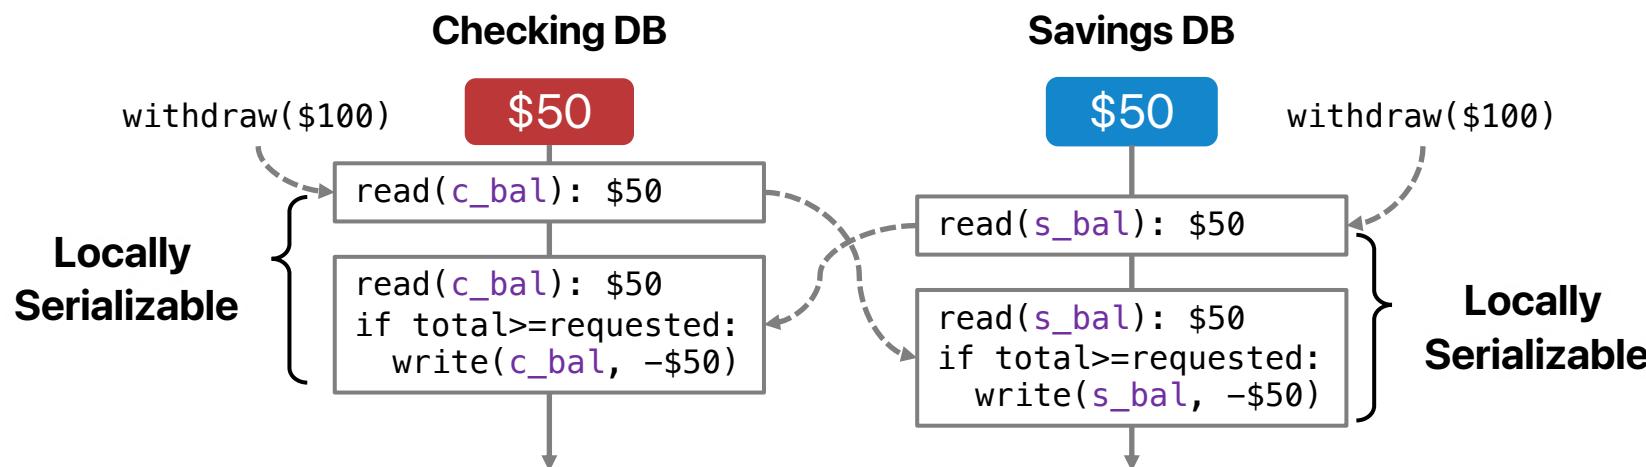
Goal: Global Serializability

- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.
- **While helpful for correctness, it is non-trivial to realize.**
 - Even if all DBs are locally serializable, global serializability can still be violated.



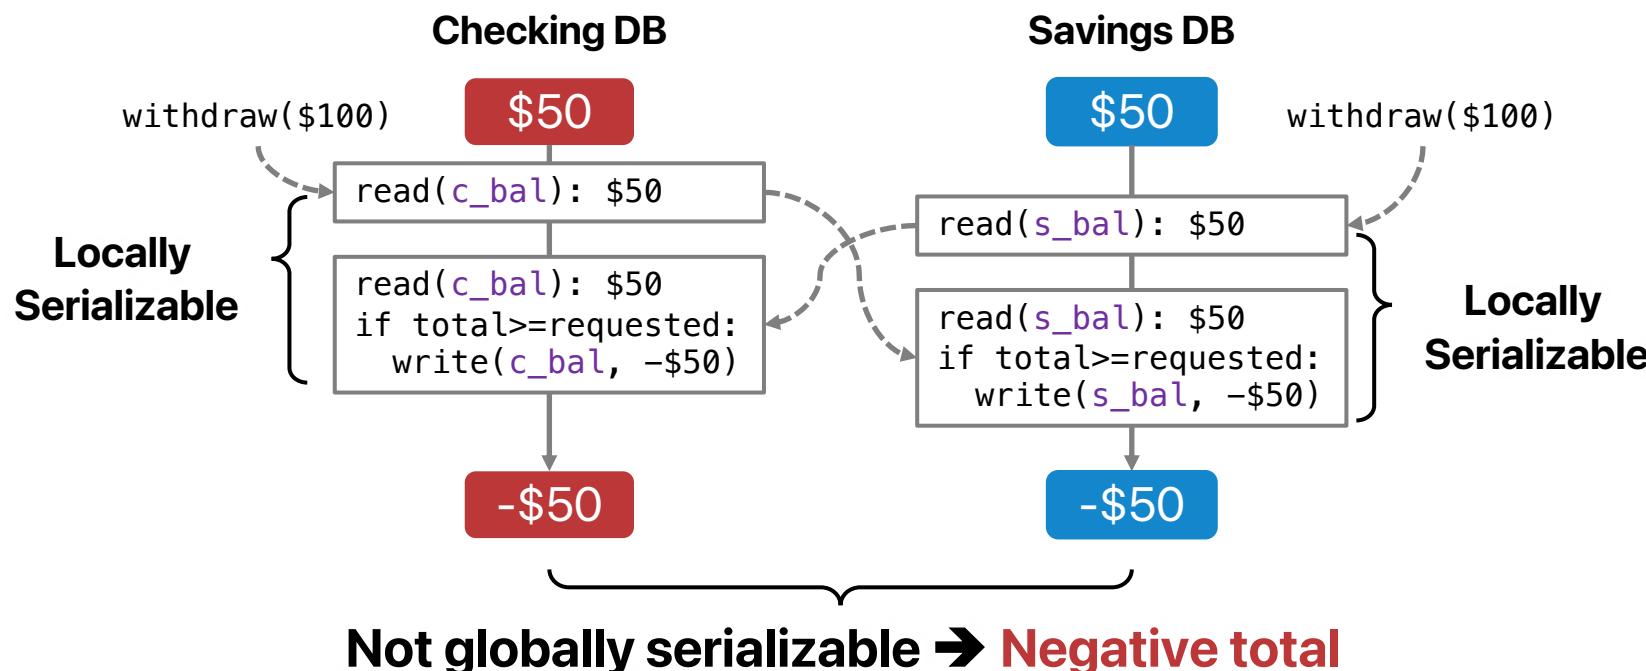
Goal: Global Serializability

- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.
- **While helpful for correctness, it is non-trivial to realize.**
 - Even if all DBs are locally serializable, global serializability can still be violated.



Goal: Global Serializability

- **Global serializability means serializable multi-DB transactions.**
 - We use "global" to distinguish from the serializability of individual local DBs.
- **While helpful for correctness, it is non-trivial to realize.**
 - Even if all DBs are locally serializable, global serializability can still be violated.



State-of-the-Art Approaches

Earlier Work [ICDE '91]*

Local DBs as serializable black boxes

- Force conflicts between subtransactions.
- **Large overhead** due to limited parallelism (>20-fold degradation)

Conservative Protocols
(e.g., forcing conflicts)



\$\$\$

Non-
Intrusive

State-of-the-Art Approaches

Earlier Work [ICDE '91]*

Local DBs as serializable black boxes

- Force conflicts between subtransactions.
- **Large overhead** due to limited parallelism (>20-fold degradation)

Conservative Protocols
(e.g., forcing conflicts)



\$\$\$

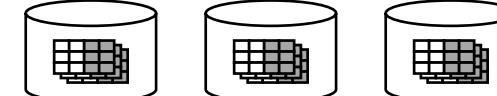
Non-
Intrusive

Recent Work [VLDB '23]**

Local DBs as non-transactional stores

- Full concurrency control in middleware.
- **Still large overhead** as DB transaction mechanisms are exercised still.
- **Intrusive changes** to for additional CC metadata.

Fully App-Level Protocols
(e.g., app-level OCC)



\$ \$

Intrusive
Metadata

* Georgakopoulos et al. 1991. On Serializability of Multidatabase Transactions Through Forced Local Conflicts. ICDE '91.

** Yamada et al. 2023. ScalarDB: Universal Transaction Manager for Polystores. VLDB '23.

State-of-the-Art Approaches

Earlier Work [ICDE '91]*

Local DBs as serializable black boxes

- Force conflicts between subtransactions.
- **Large overhead** due to limited parallelism (>20-fold degradation)

Conservative Protocols
(e.g., forcing conflicts)

\$\$\$



Non-
Intrusive

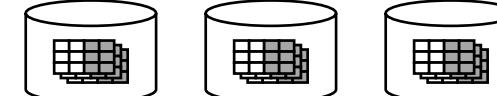
Recent Work [VLDB '23]**

Local DBs as non-transactional stores

- Full concurrency control in middleware.
- **Still large overhead** as DB transaction mechanisms are exercised still.
- **Intrusive changes** to for additional CC metadata.

Fully App-Level Protocols
(e.g., app-level OCC)

\$ \$



Intrusive
Metadata

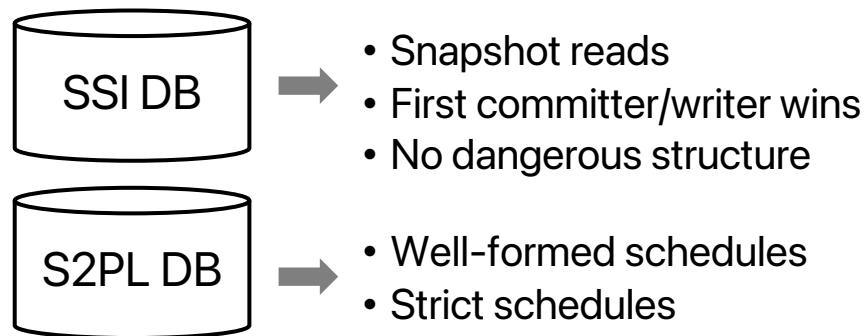
**Can we achieve high performance
without intrusive changes?**

* Georgakopoulos et al. 1991. On Serializability of Multidatabase Transactions Through Forced Local Conflicts. ICDE '91.

** Yamada et al. 2023. ScalarDB: Universal Transaction Manager for Polystores. VLDB '23.

Our Gray-Box Approach

- **Observation:** Dominant CC families in mainstream DBs, SSI & S2PL, offer useful common properties.



Our Gray-Box Approach

- **Observation:** Dominant CC families in mainstream DBs, SSI & S2PL, offer useful common properties.
- **Basic Idea:** Reuse such properties & add necessary coordination.



Our Gray-Box Approach

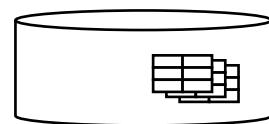
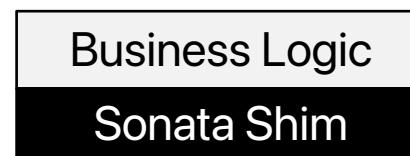
- **Observation:** Dominant CC families in mainstream DBs, SSI & S2PL, offer useful common properties.
- **Basic Idea:** Reuse such properties & add necessary coordination.



The theory of commitment ordering (CO) [VLDB '92] enables a locally enforceable condition, allowing lightweight coordination.

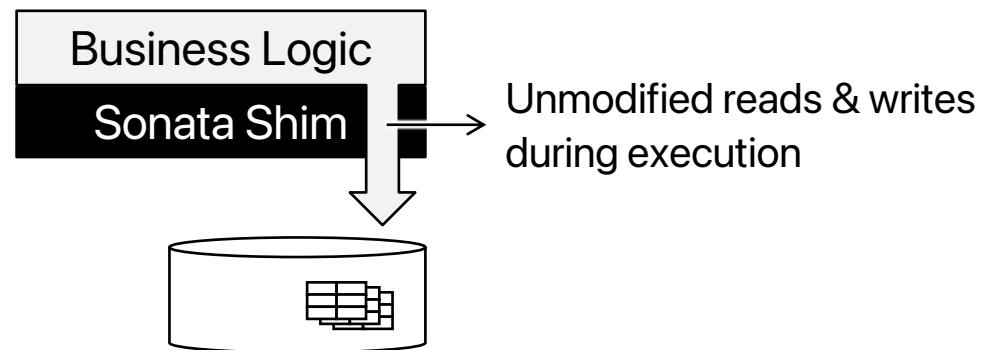
Sonata Overview

- Sonata works as application-level shims.



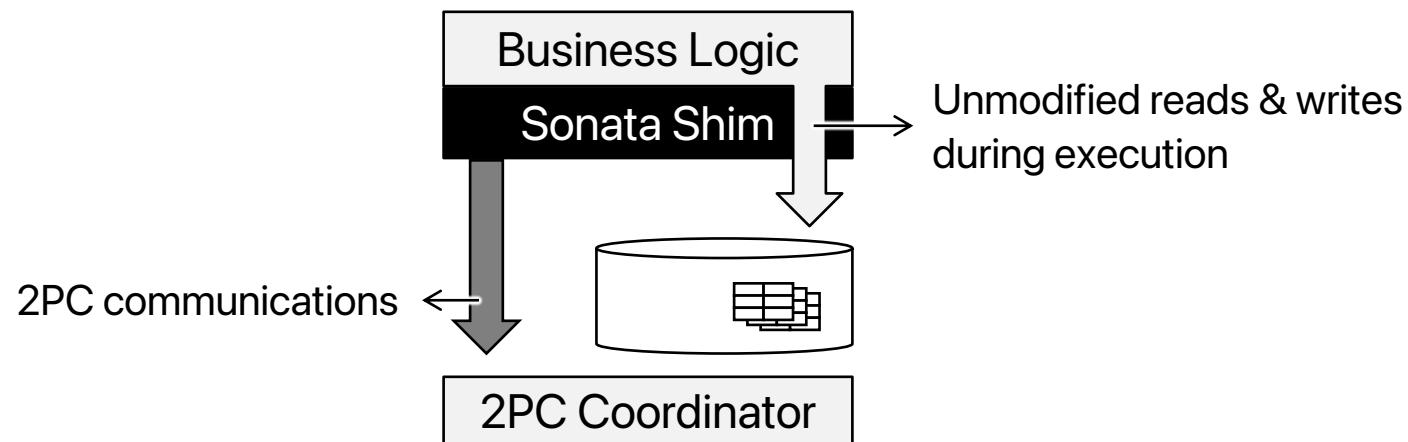
Sonata Overview

- **Sonata works as application-level shims.**
 - **No change** to apps' schemas, query statements, the DB drivers, or the DB systems.



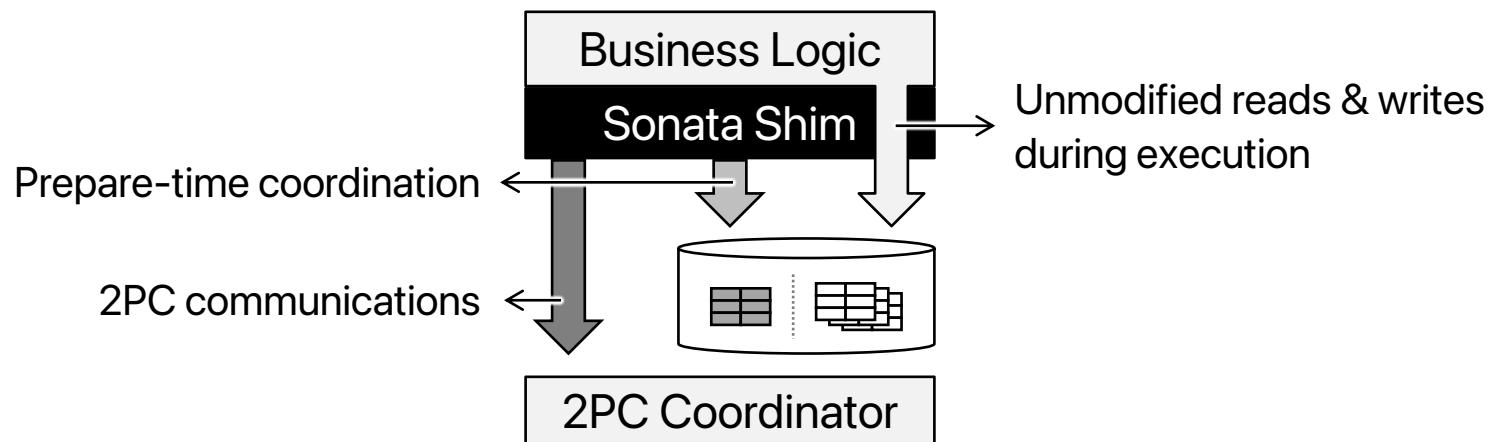
Sonata Overview

- **Sonata works as application-level shims.**
 - **No change** to apps' schemas, query statements, the DB drivers, or the DB systems.
 - Use 2PC for atomicity and durability.



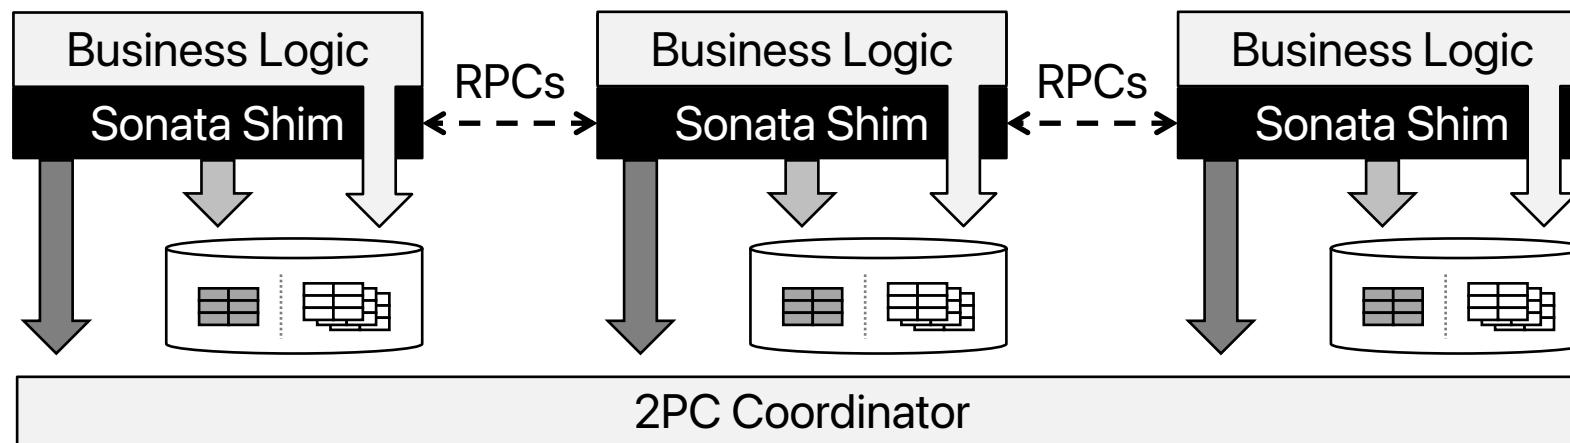
Sonata Overview

- **Sonata works as application-level shims.**
 - **No change** to apps' schemas, query statements, the DB drivers, or the DB systems.
 - Use 2PC for atomicity and durability.
 - Take actions **only at 2PC prepare time** independently at individual DBs.



Sonata Overview

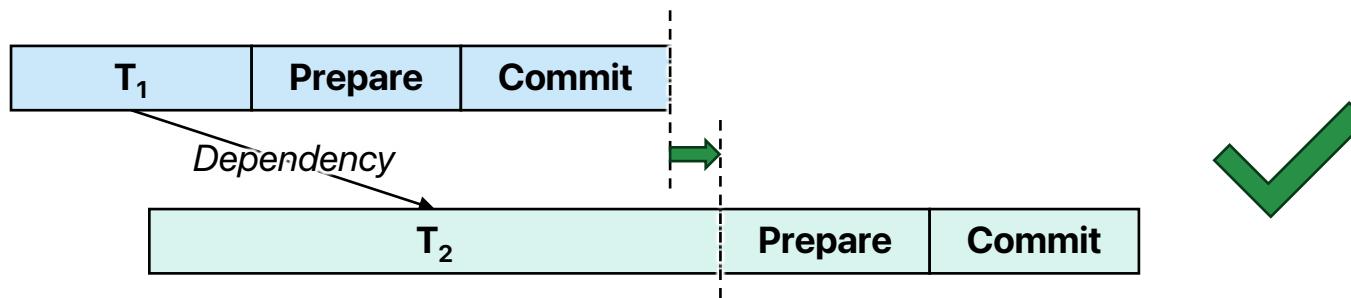
- **Sonata works as application-level shims.**
 - **No change** to apps' schemas, query statements, the DB drivers, or the DB systems.
 - Use 2PC for atomicity and durability.
 - Take actions **only at 2PC prepare time** independently at individual DBs.



(Also works with single-server, multi-DB cases)

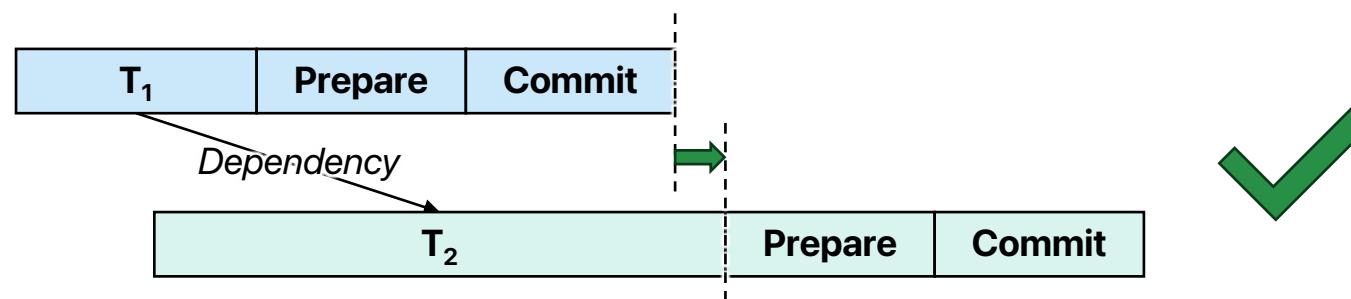
Global Serializability via Local Enforcement

- **The local condition** (derived from the CO theory):
In a DB, if $T_1 \rightarrow T_2$, then T_1 commits before T_2 prepares.



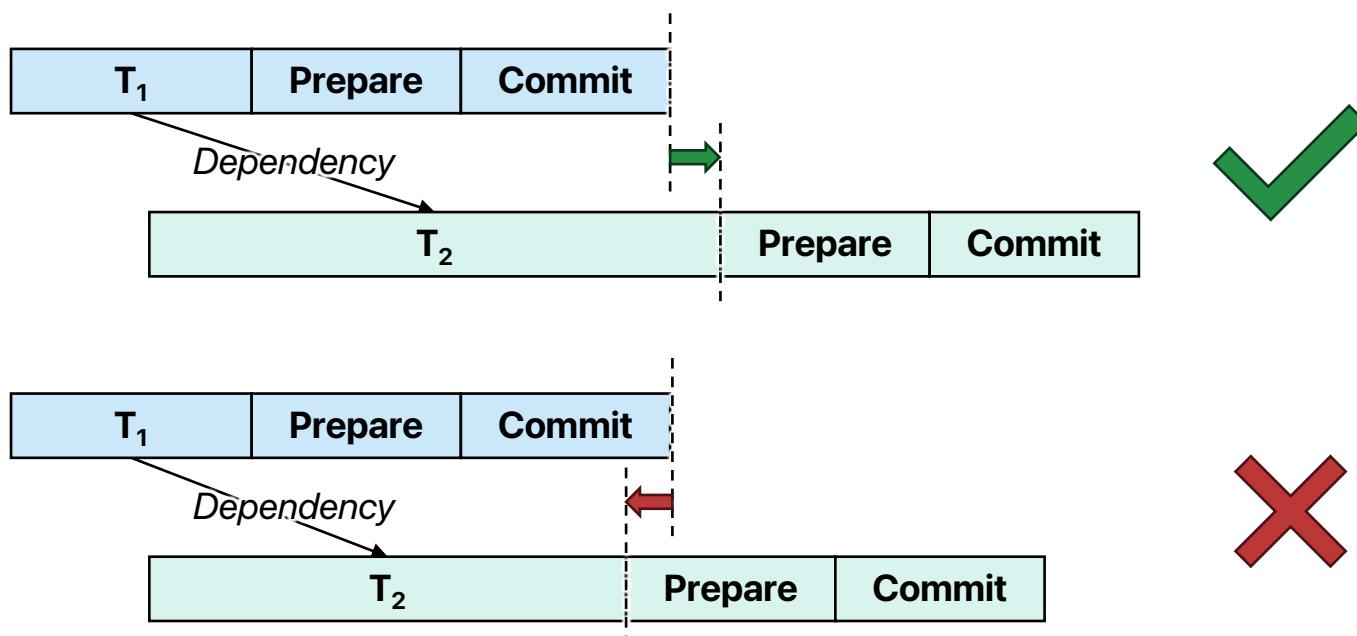
Global Serializability via Local Enforcement

- **The local condition** (derived from the CO theory):
In a DB, if $T_1 \rightarrow T_2$, then T_1 commits before T_2 prepares.
- **If all DBs satisfies this condition, global serializability holds.**



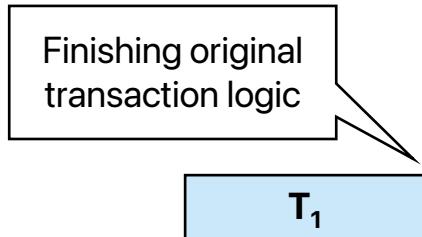
Global Serializability via Local Enforcement

- **The local condition** (derived from the CO theory):
In a DB, if $T_1 \rightarrow T_2$, then T_1 commits before T_2 prepares.
- If all DBs satisfies this condition, global serializability holds.



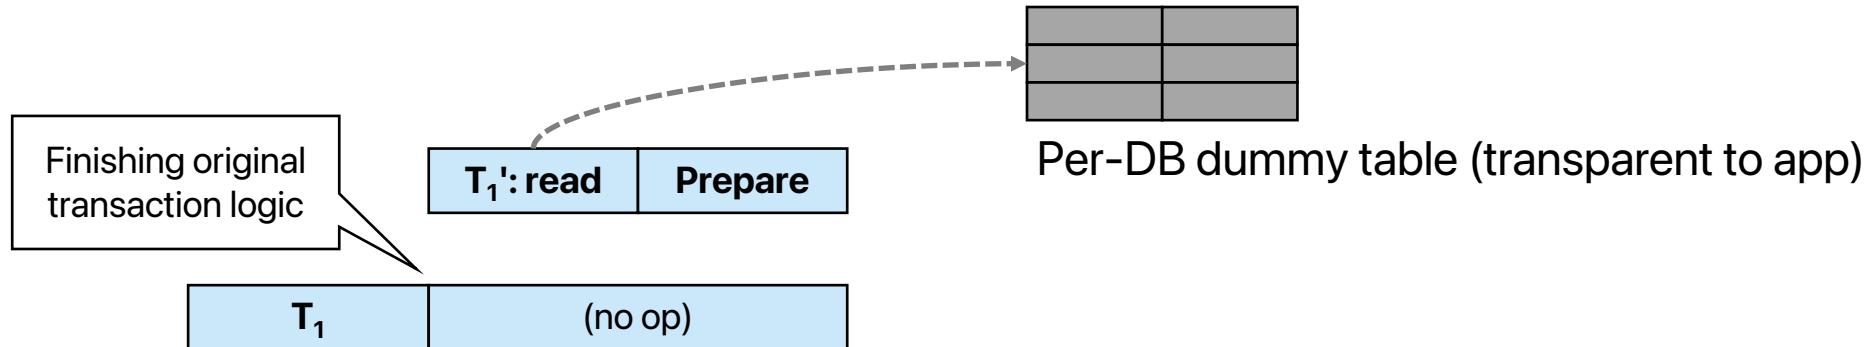
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



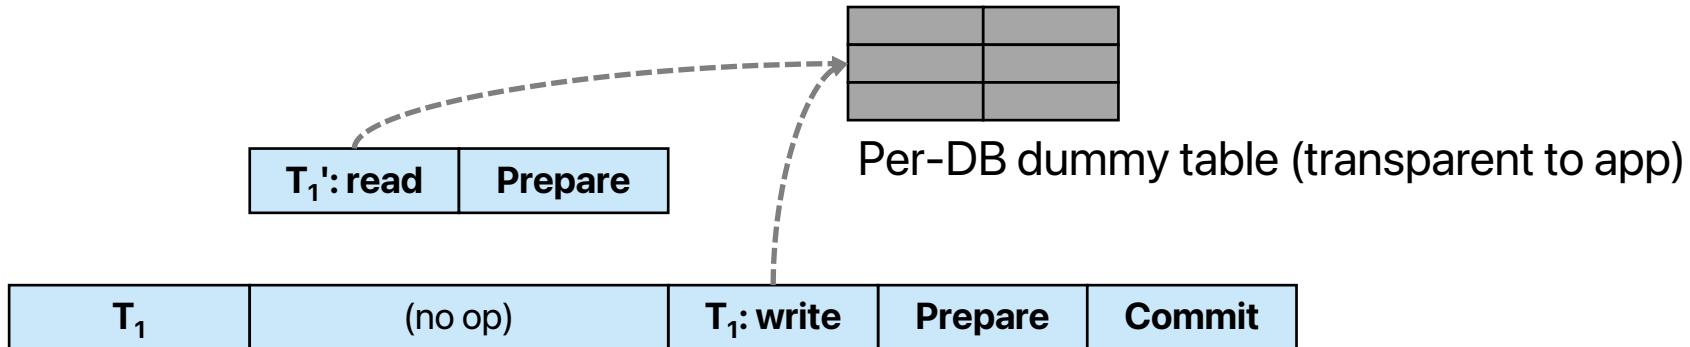
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



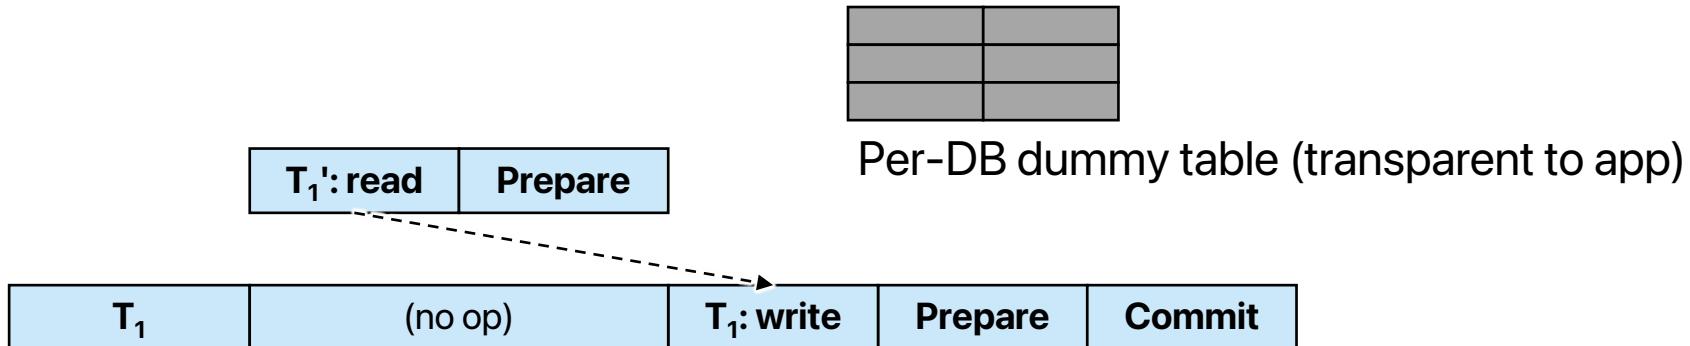
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



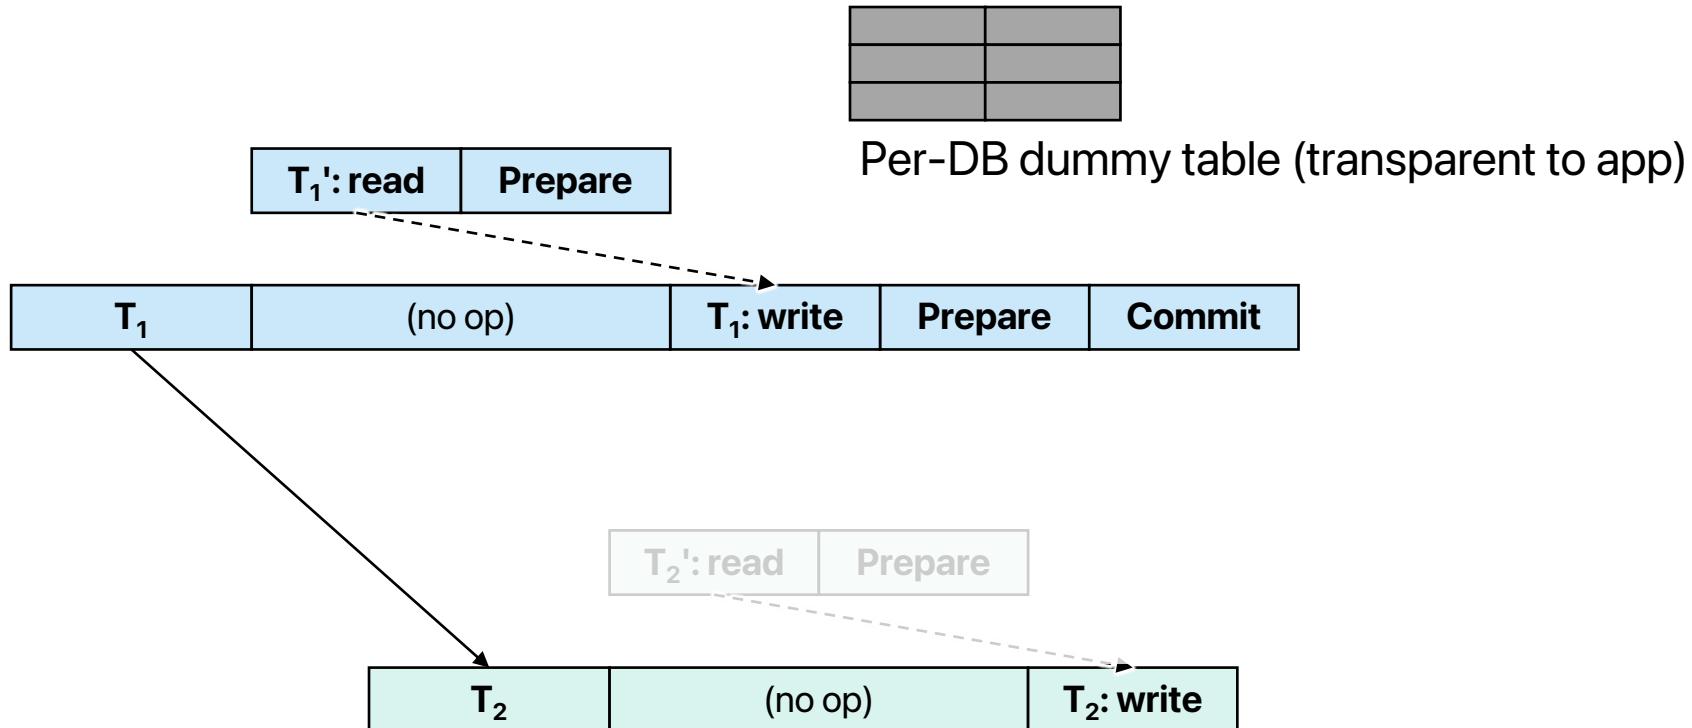
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



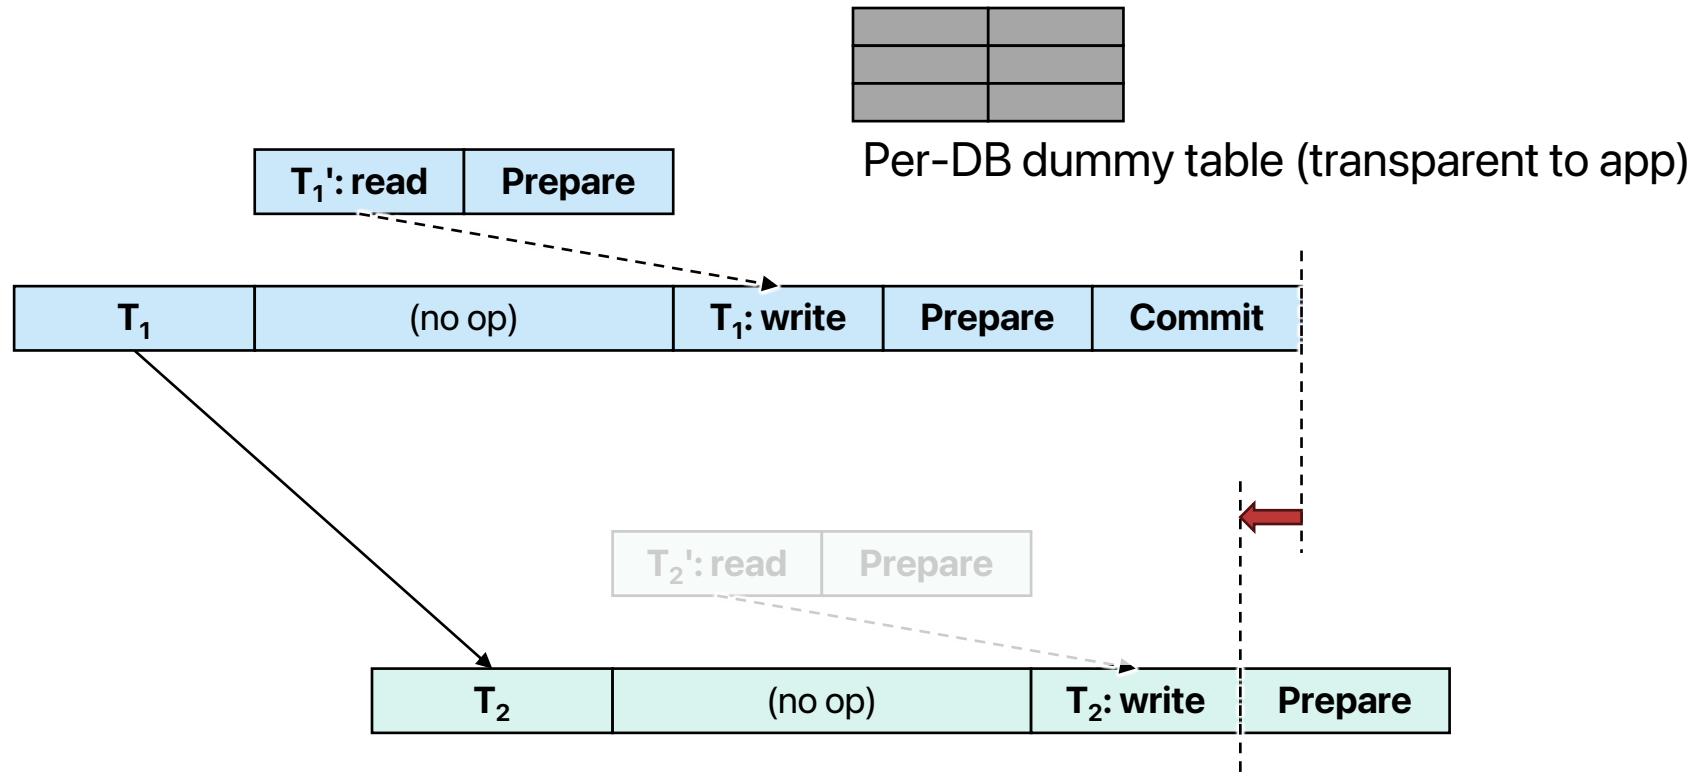
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



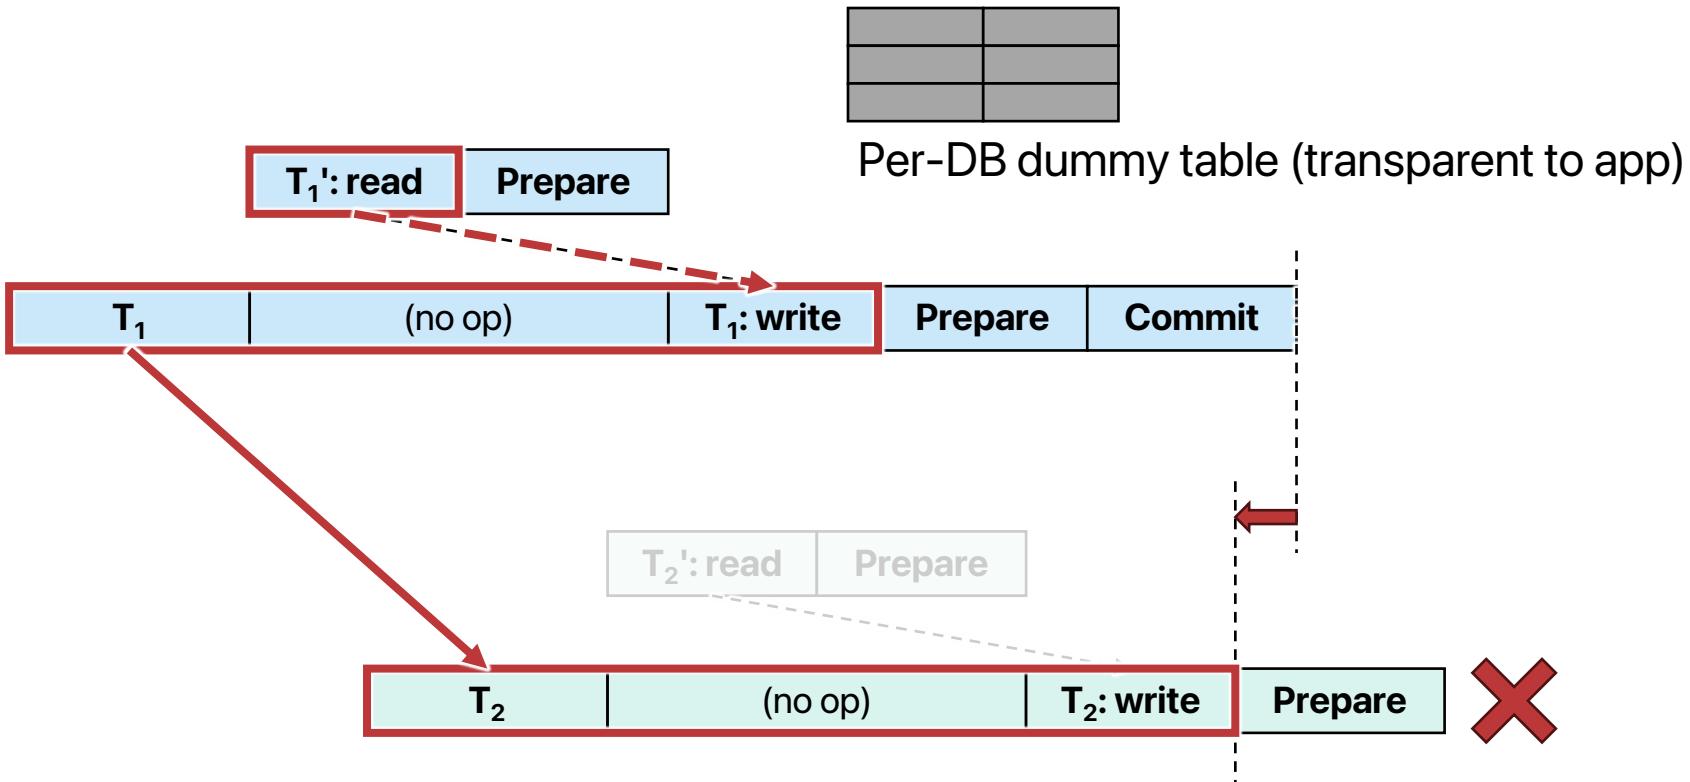
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



Enforcement in SSI DBs

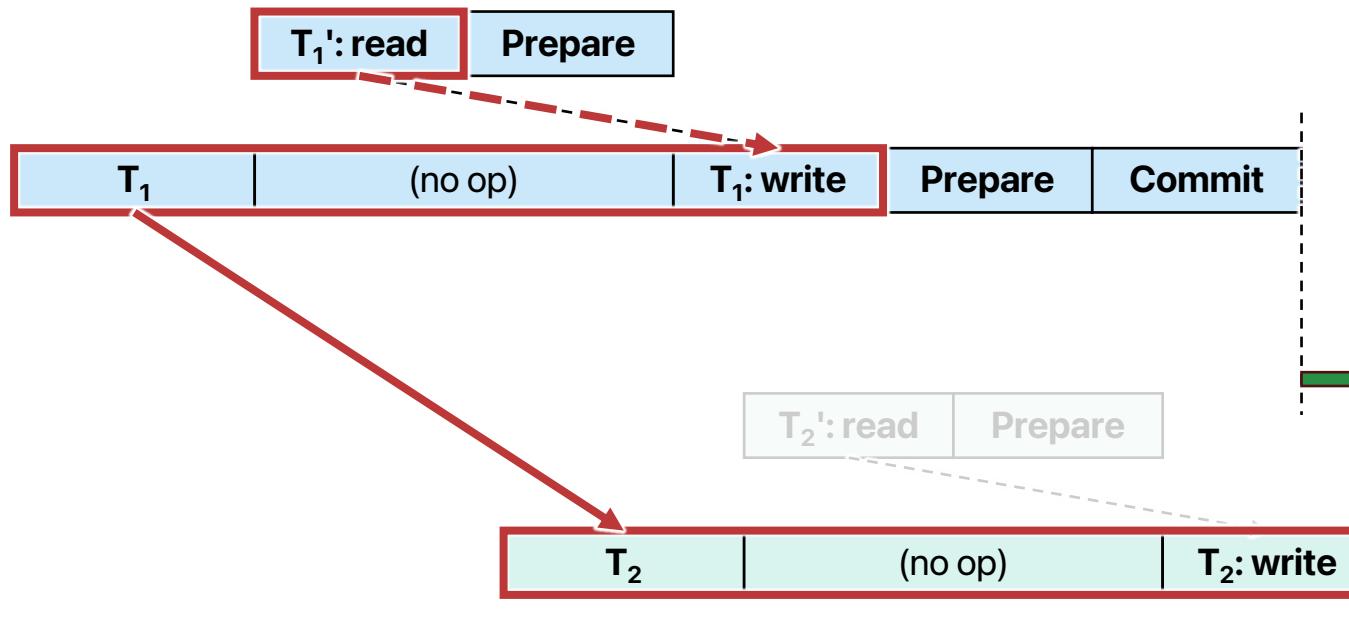
- For SSI DBs, we introduce helper transactions.



SSI dangerous structure!
T₂ will be aborted → No violation

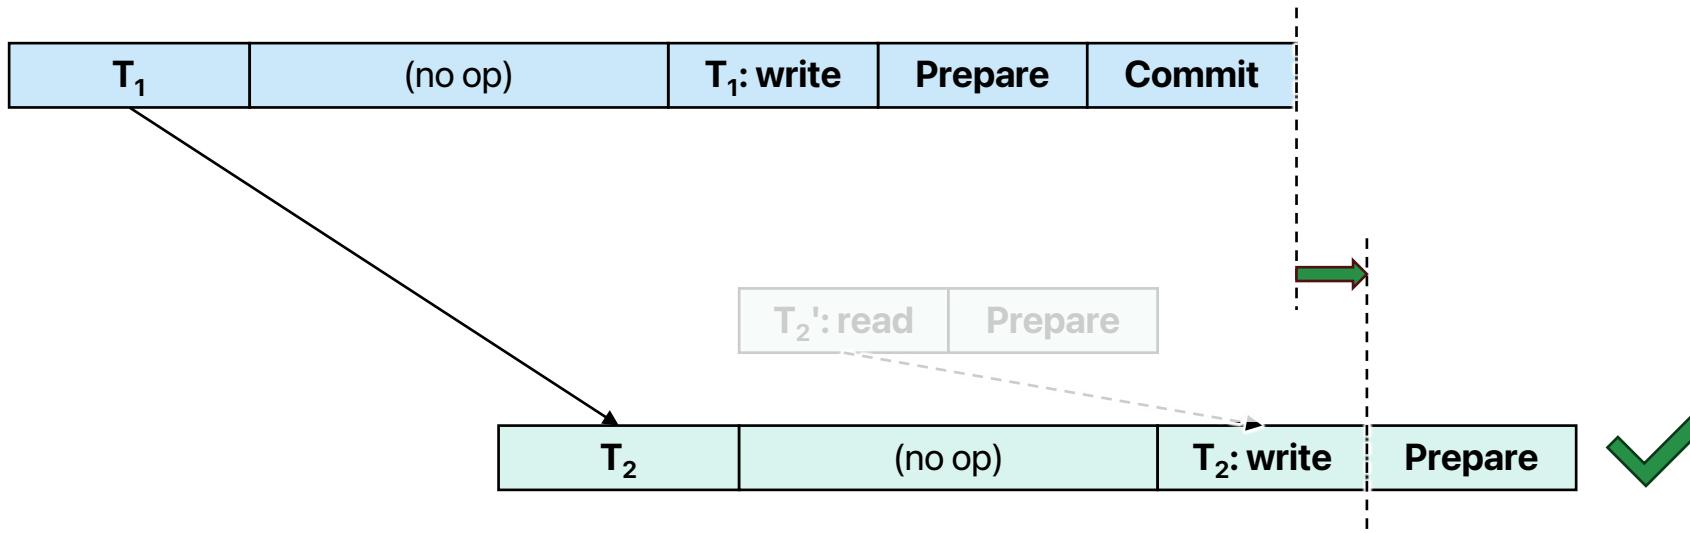
Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



Enforcement in SSI DBs

- For SSI DBs, we introduce helper transactions.



No SSI dangerous structure!
T2 won't be aborted → No false positive

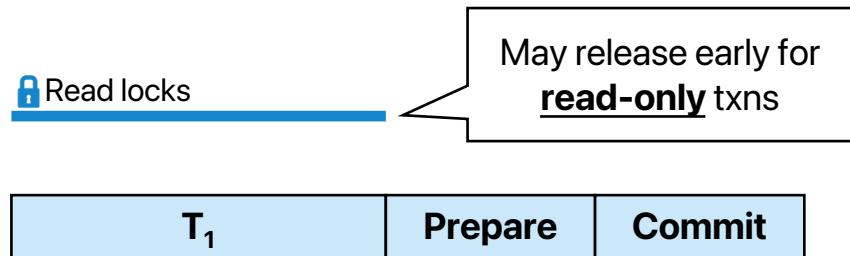
Enforcement in S2PL DBs

- For S2PL DBs, we simply introduce dummy writes.



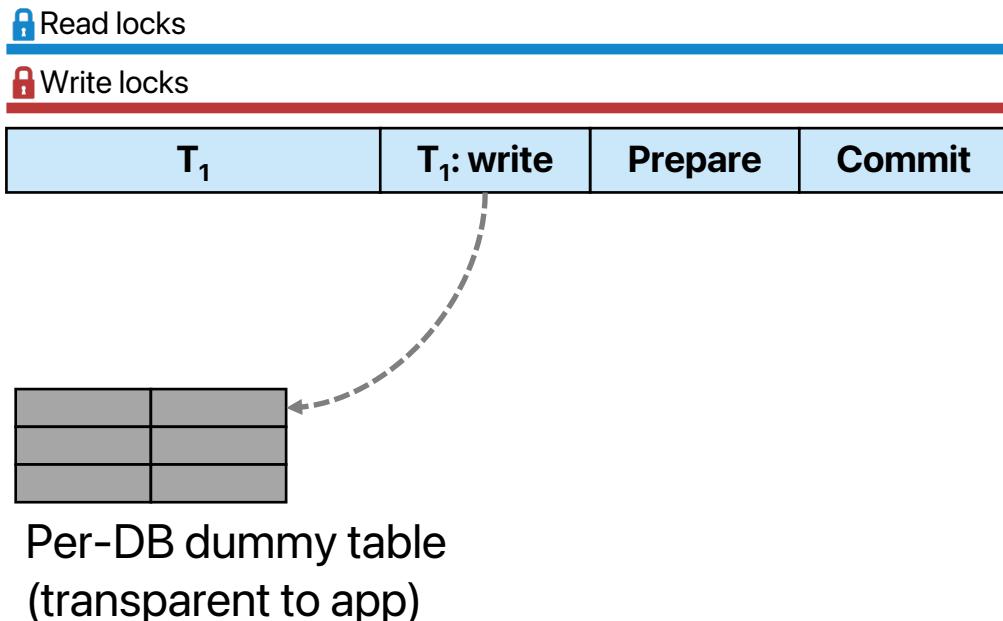
Enforcement in S2PL DBs

- For S2PL DBs, we simply introduce dummy writes.



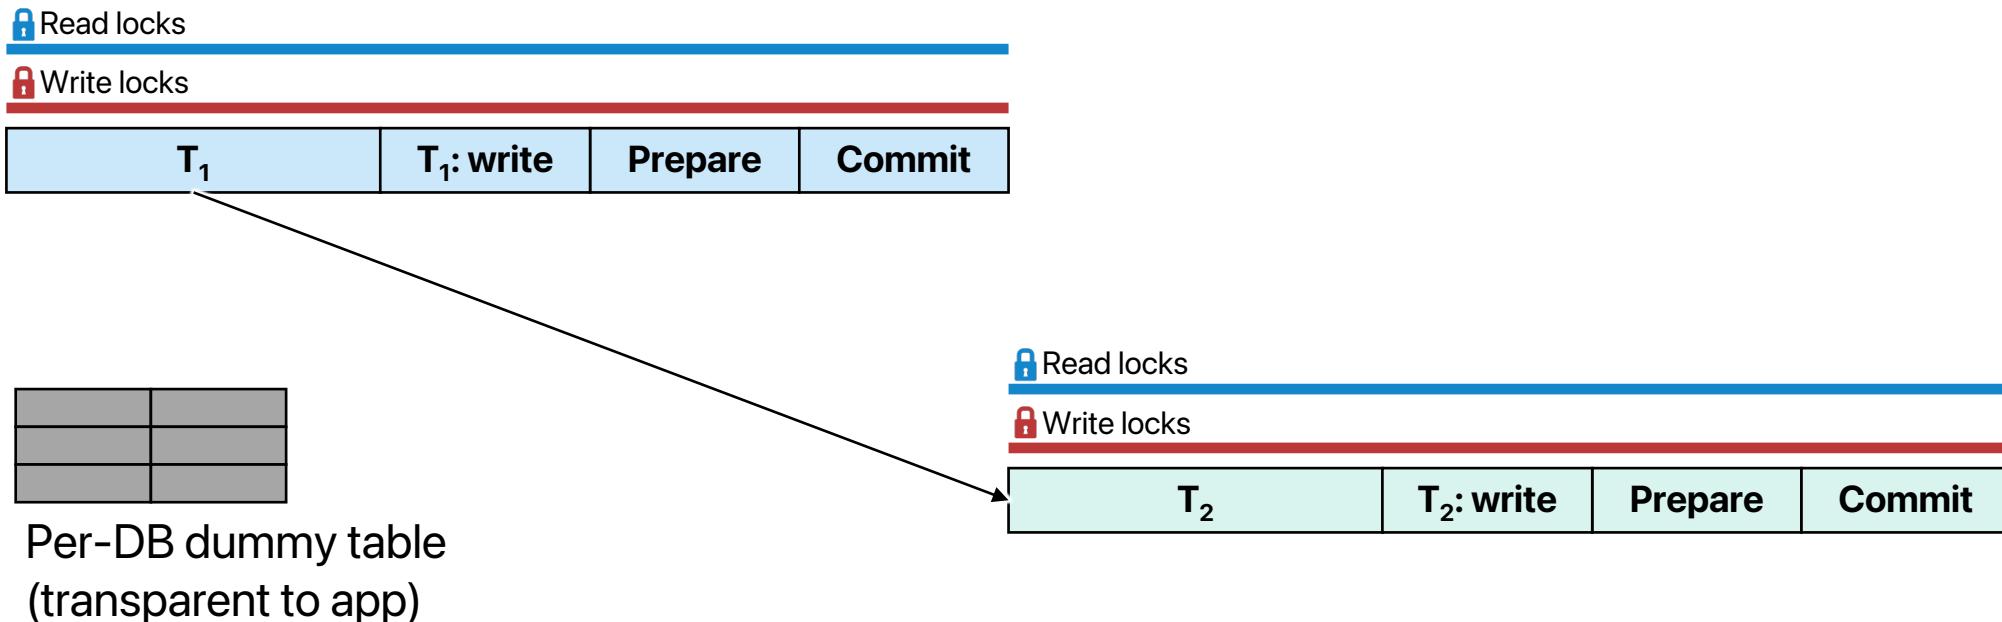
Enforcement in S2PL DBs

- For S2PL DBs, we simply introduce dummy writes.



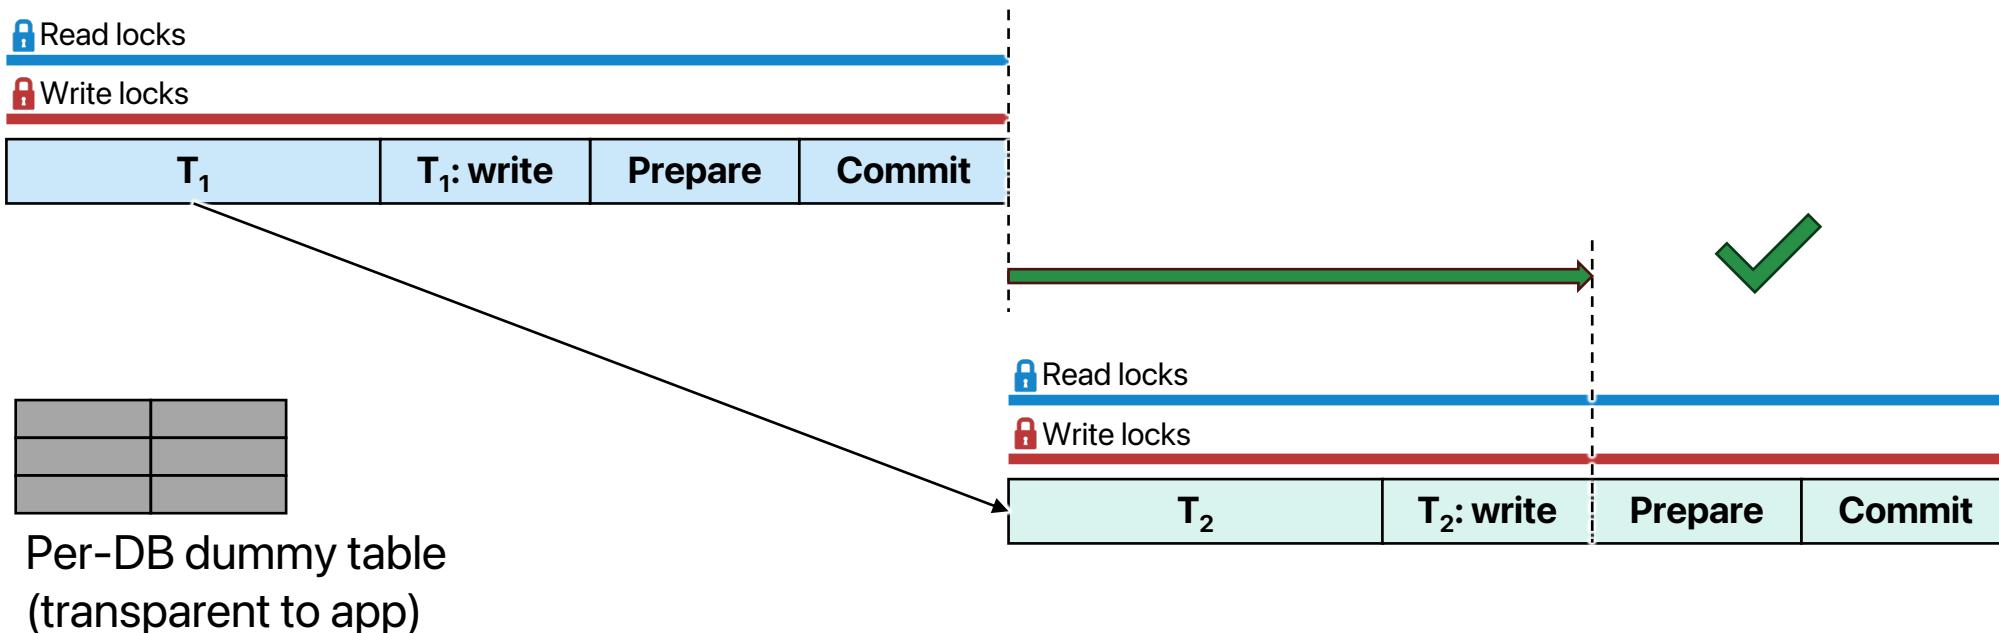
Enforcement in S2PL DBs

- For S2PL DBs, we simply introduce dummy writes.



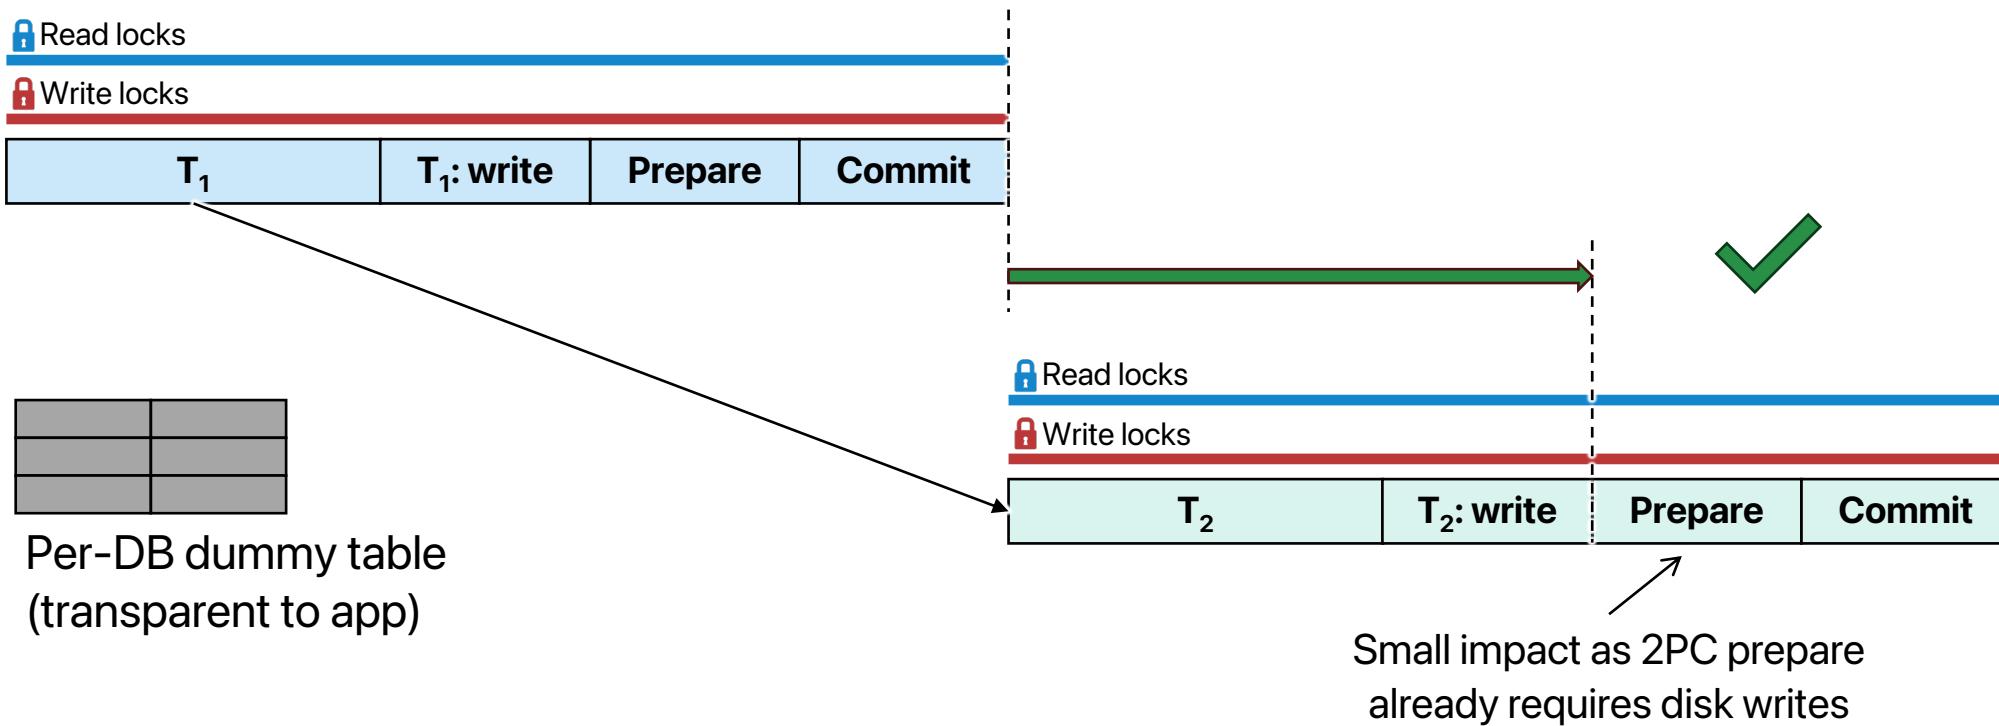
Enforcement in S2PL DBs

- For S2PL DBs, we simply introduce dummy writes.



Enforcement in S2PL DBs

- For S2PL DBs, we simply introduce dummy writes.



Correctness and False Positives

- **Sonata provably guarantees the local condition, thus global serializability.** (Sec. 4.1)
 - The proofs leverage lemmas from [TODS 05]*.
- **With PostgreSQL and MySQL, there's no false positive if the DB's conflict detection is accurate.** (Sec. 4.2)
 - The analysis depends on DB's implementation details.

Programming with Sonata

- Our prototype is in Java and is based on Apache Seata (25k+ GH ★).

Programming with Sonata

- Our prototype is in Java and is based on Apache Seata (25k+ GH ★).
- For existing applications, codebase changes can be as small as adding a few annotations.

```
@GlobalTransactional
public void workflow() {
    // Access the local database
    jdbcClient.sql("Select ...").query();

    // Request other services
    restClient.post().uri( uri: "/api" ).retrieve().toBodilessEntity();
}
```

Programming with Sonata

- Our prototype is in Java and is based on Apache Seata (25k+ GH ★).
- For existing applications, codebase changes can be as small as adding a few annotations.

```
@GlobalTransactional  
public void workflow() {  
    // Access the local database  
    jdbcClient.sql("Select ...").query(); → JDBC or any JDBC-based library  
                                                (We intercept begin/commit calls)  
  
    // Request other services  
    restClient.post().uri( uri: "/api" ).retrieve().toBodilessEntity();  
}
```

Programming with Sonata

- Our prototype is in Java and is based on Apache Seata (25k+ GH ★).
- For existing applications, codebase changes can be as small as adding a few annotations.

```
@GlobalTransactional  
public void workflow() {  
    // Access the local database  
    jdbcClient.sql("Select ...").query(); → JDBC or any JDBC-based library  
                                                (We intercept begin/commit calls)  
  
    // Request other services  
    restClient.post().uri( uri: "/api").retrieve().toBodilessEntity();  
}
```

→ Any RPC framework able to piggyback transaction ID strings
(We currently intercept HTTP requests)

Sonata Performance

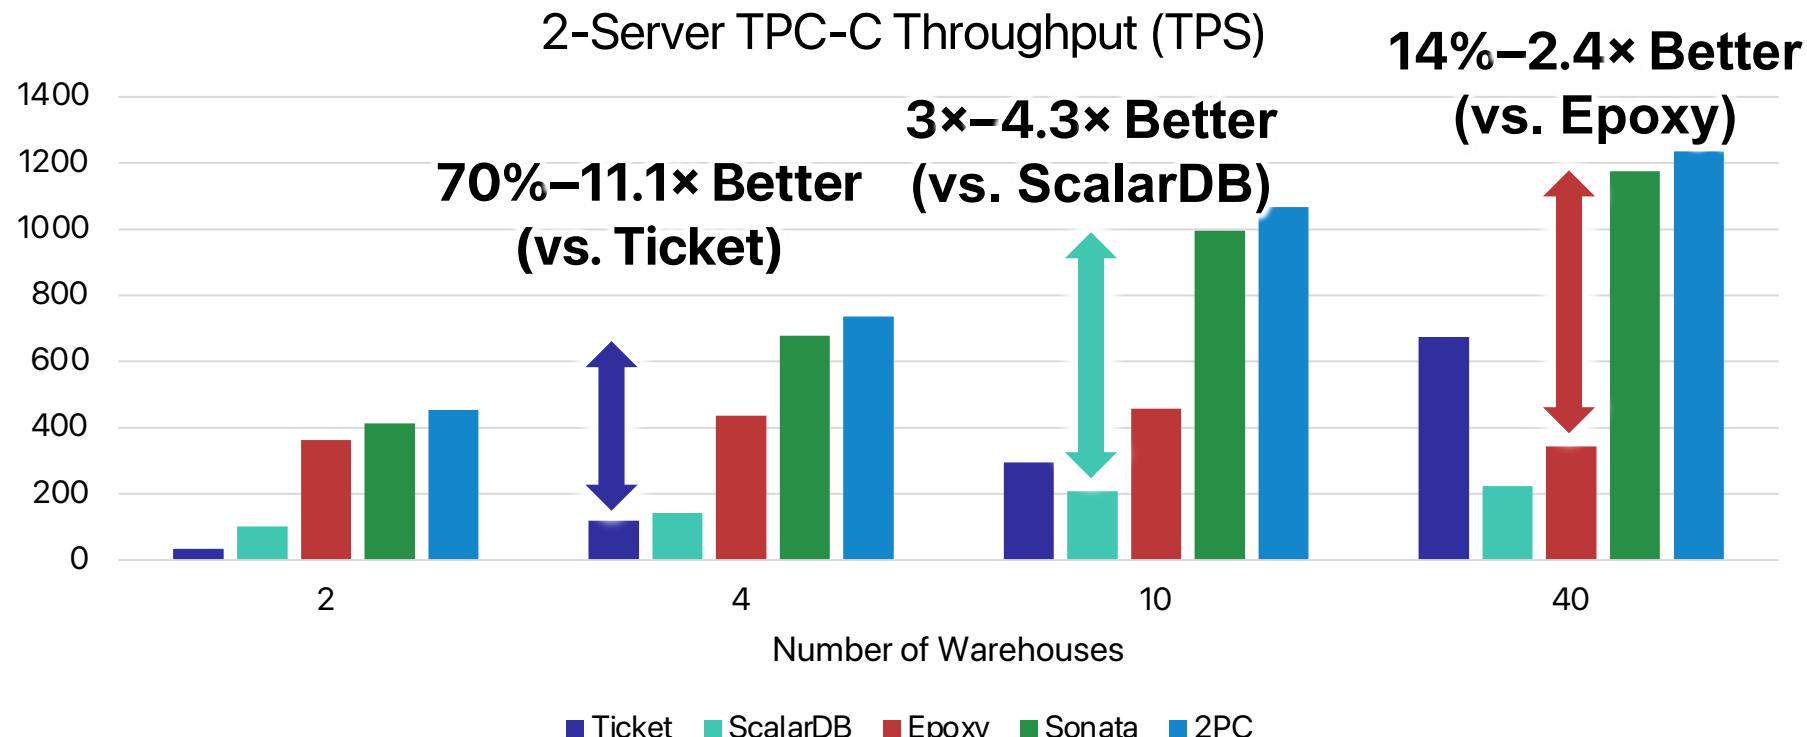
- We build a multi-DB TPC-C by partitioning by warehouses.
 - Either PostgreSQL or MySQL as local DBs.
 - Global serializability: ScalarDB [VLDB '23] & Ticket [ICDE '91].

Sonata Performance

- We build a multi-DB TPC-C by partitioning by warehouses.
 - Either PostgreSQL or MySQL as local DBs.
 - Global serializability: ScalarDB [VLDB '23] & Ticket [ICDE '91].
 - Snapshot isolation: Epoxy* [VLDB '23].

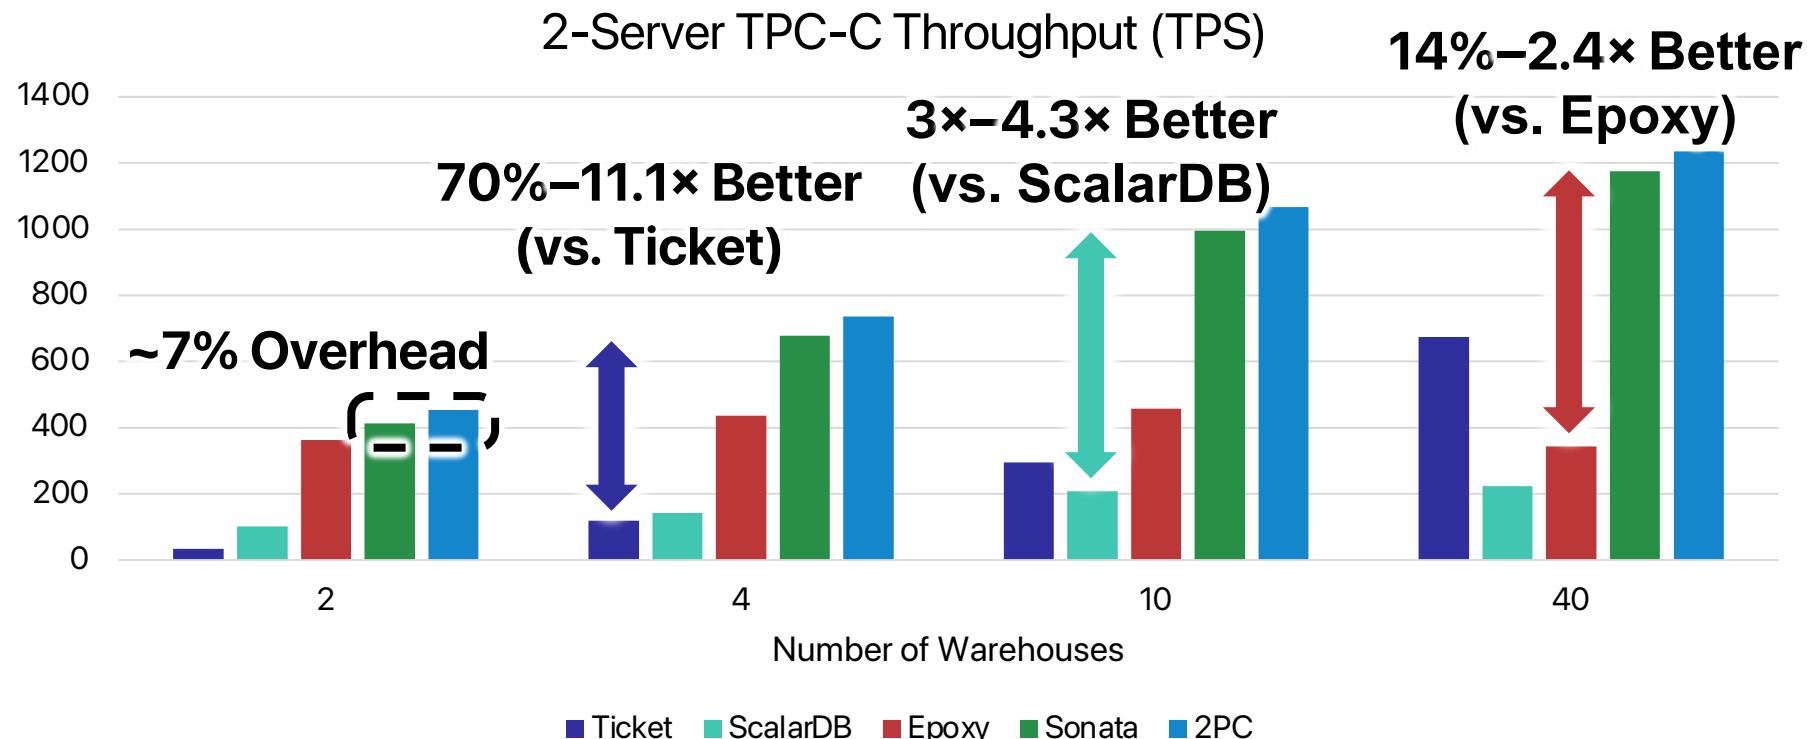
Sonata Performance

- We build a multi-DB TPC-C by partitioning by warehouses.
 - Either PostgreSQL or MySQL as local DBs.
 - Global serializability: ScalarDB [VLDB '23] & Ticket [ICDE '91].
 - Snapshot isolation: Epoxy* [VLDB '23].



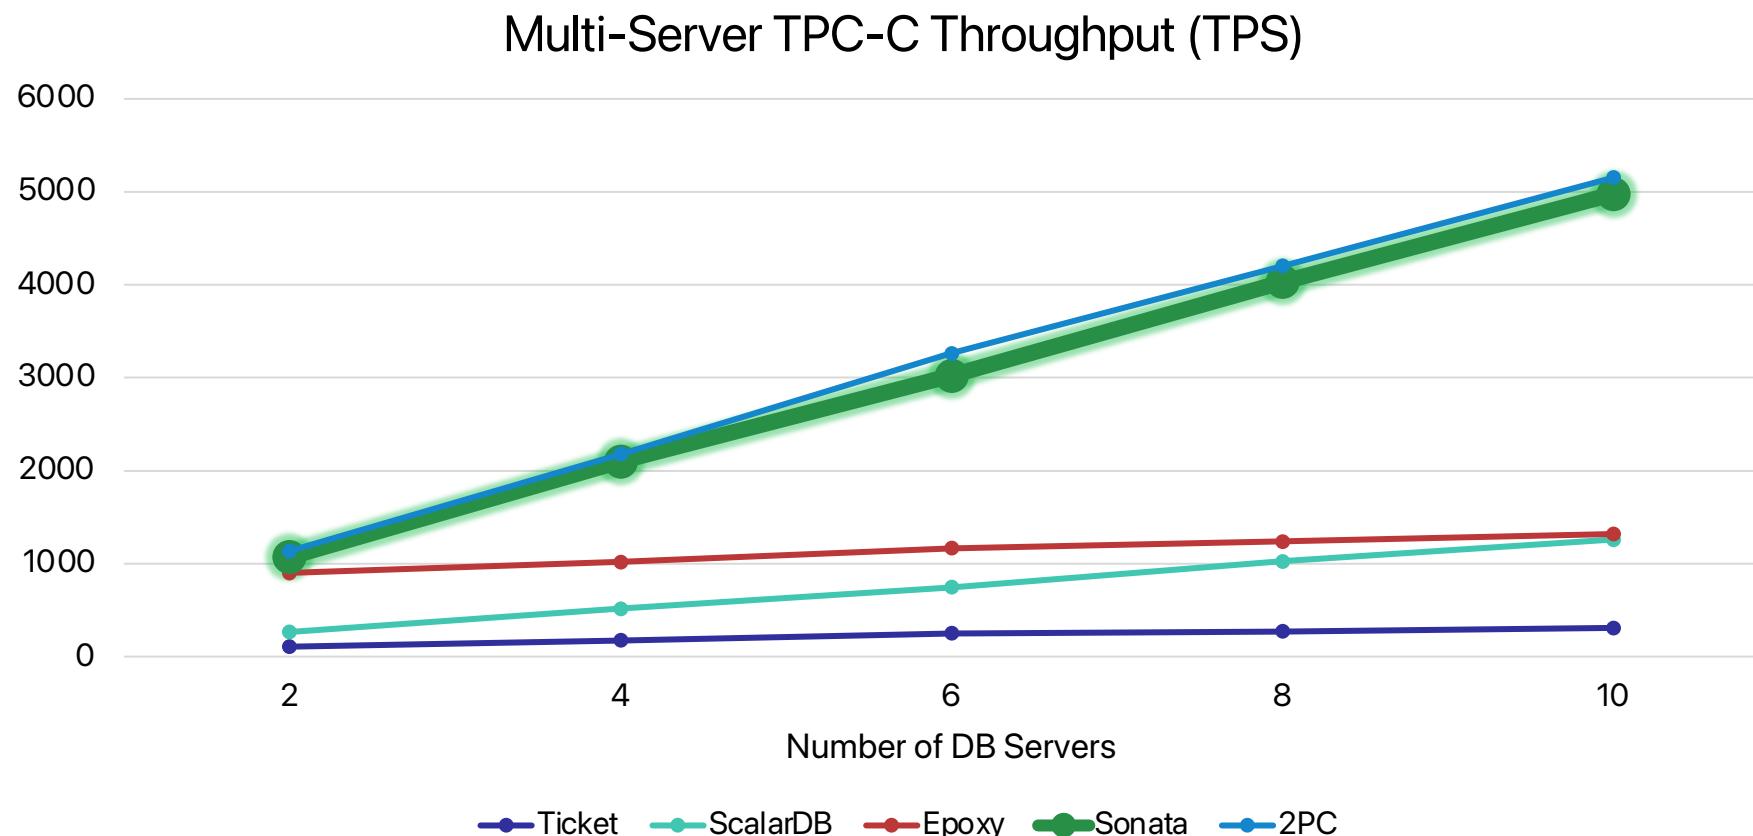
Sonata Performance

- We build a multi-DB TPC-C by partitioning by warehouses.
 - Either PostgreSQL or MySQL as local DBs.
 - Global serializability: ScalarDB [VLDB '23] & Ticket [ICDE '91].
 - Snapshot isolation: Epoxy* [VLDB '23].



Sonata Performance

- For scalability, at 10 nodes, Sonata (green) achieves $4.65 \times$ its 2-node throughput.



Summary and Q/A

- Sonata offers global serializability in an **efficient, non-intrusive way**.
- It **reuses DB concurrency control** to add **lightweight coordination**.
- It achieves up to **11.1× higher throughput** than prior solutions.

- Check out our paper and code!
 - Design details, optimizations, proofs, & more evaluation results.



Paper



Code